

Developing a Network Virtualization Framework for Testing Network Resilience Techniques

Undergraduate Thesis

Pablo Collado Soto - 29/06/2021


[Source](#)



[Source](#)

Project's Background

- Testing environment for the REACT framework.
- REACT -> allows the reconfiguration of networks to mitigate attacks.
- Testing environment -> fully virtual network.
- Reconfiguration -> altering network topology “on the fly”.
- Flexibility -> “in-house” solution.



[Source](#)

Initial Technology Choices

- *Layer 3 (L3)* -> collection of hosts connected by routers.

- Hosts and routers -> virtual entities.

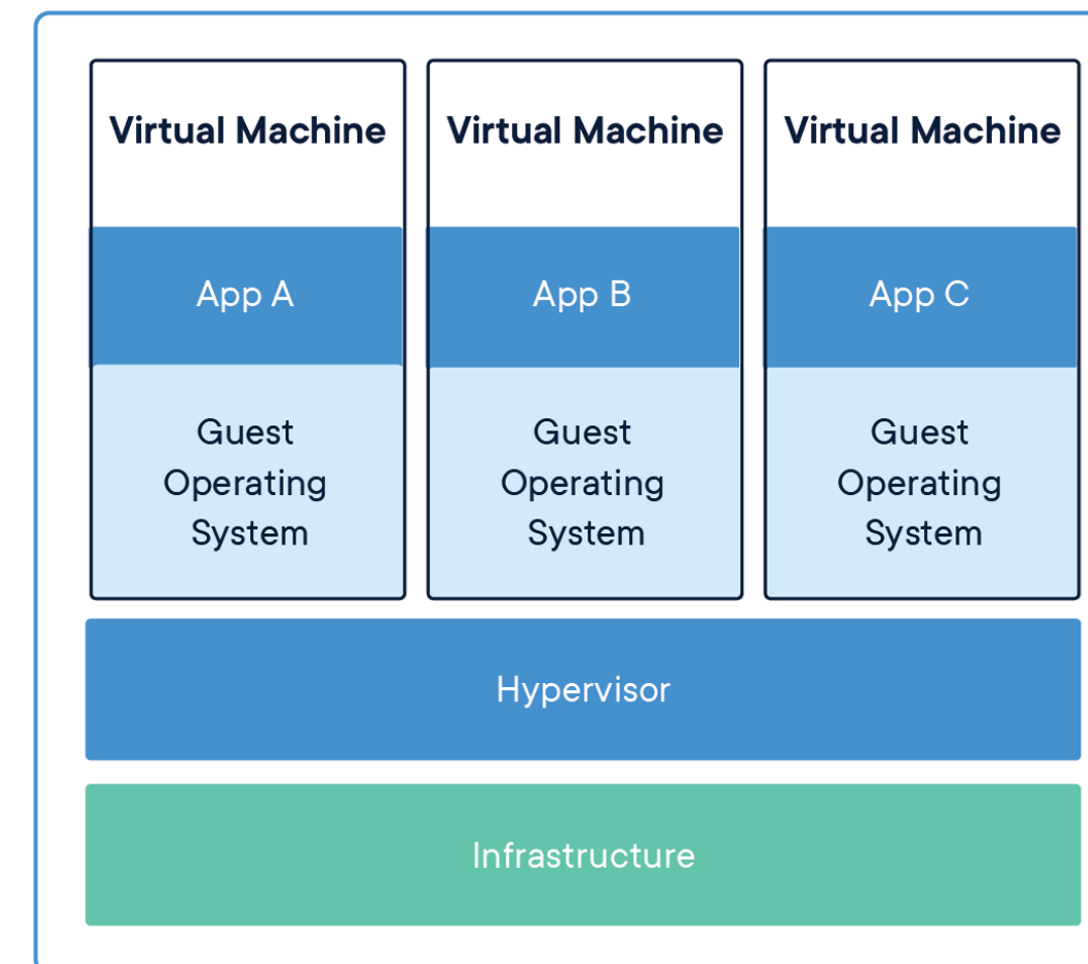
- Two contenders:

- **Virtual Machines**

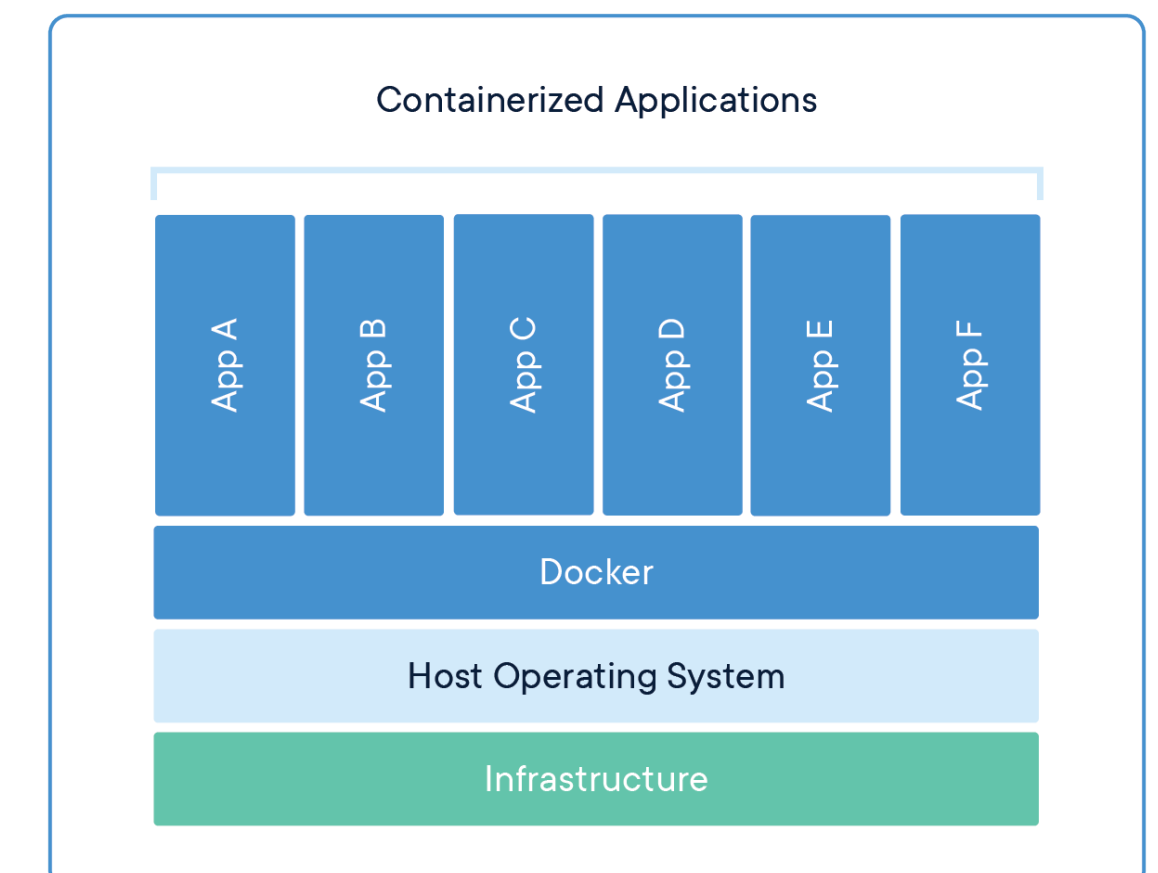
- Resource intensive.
- “Obscure” network configuration.

- **Containers**

- Lightweight.
- Finer network control.



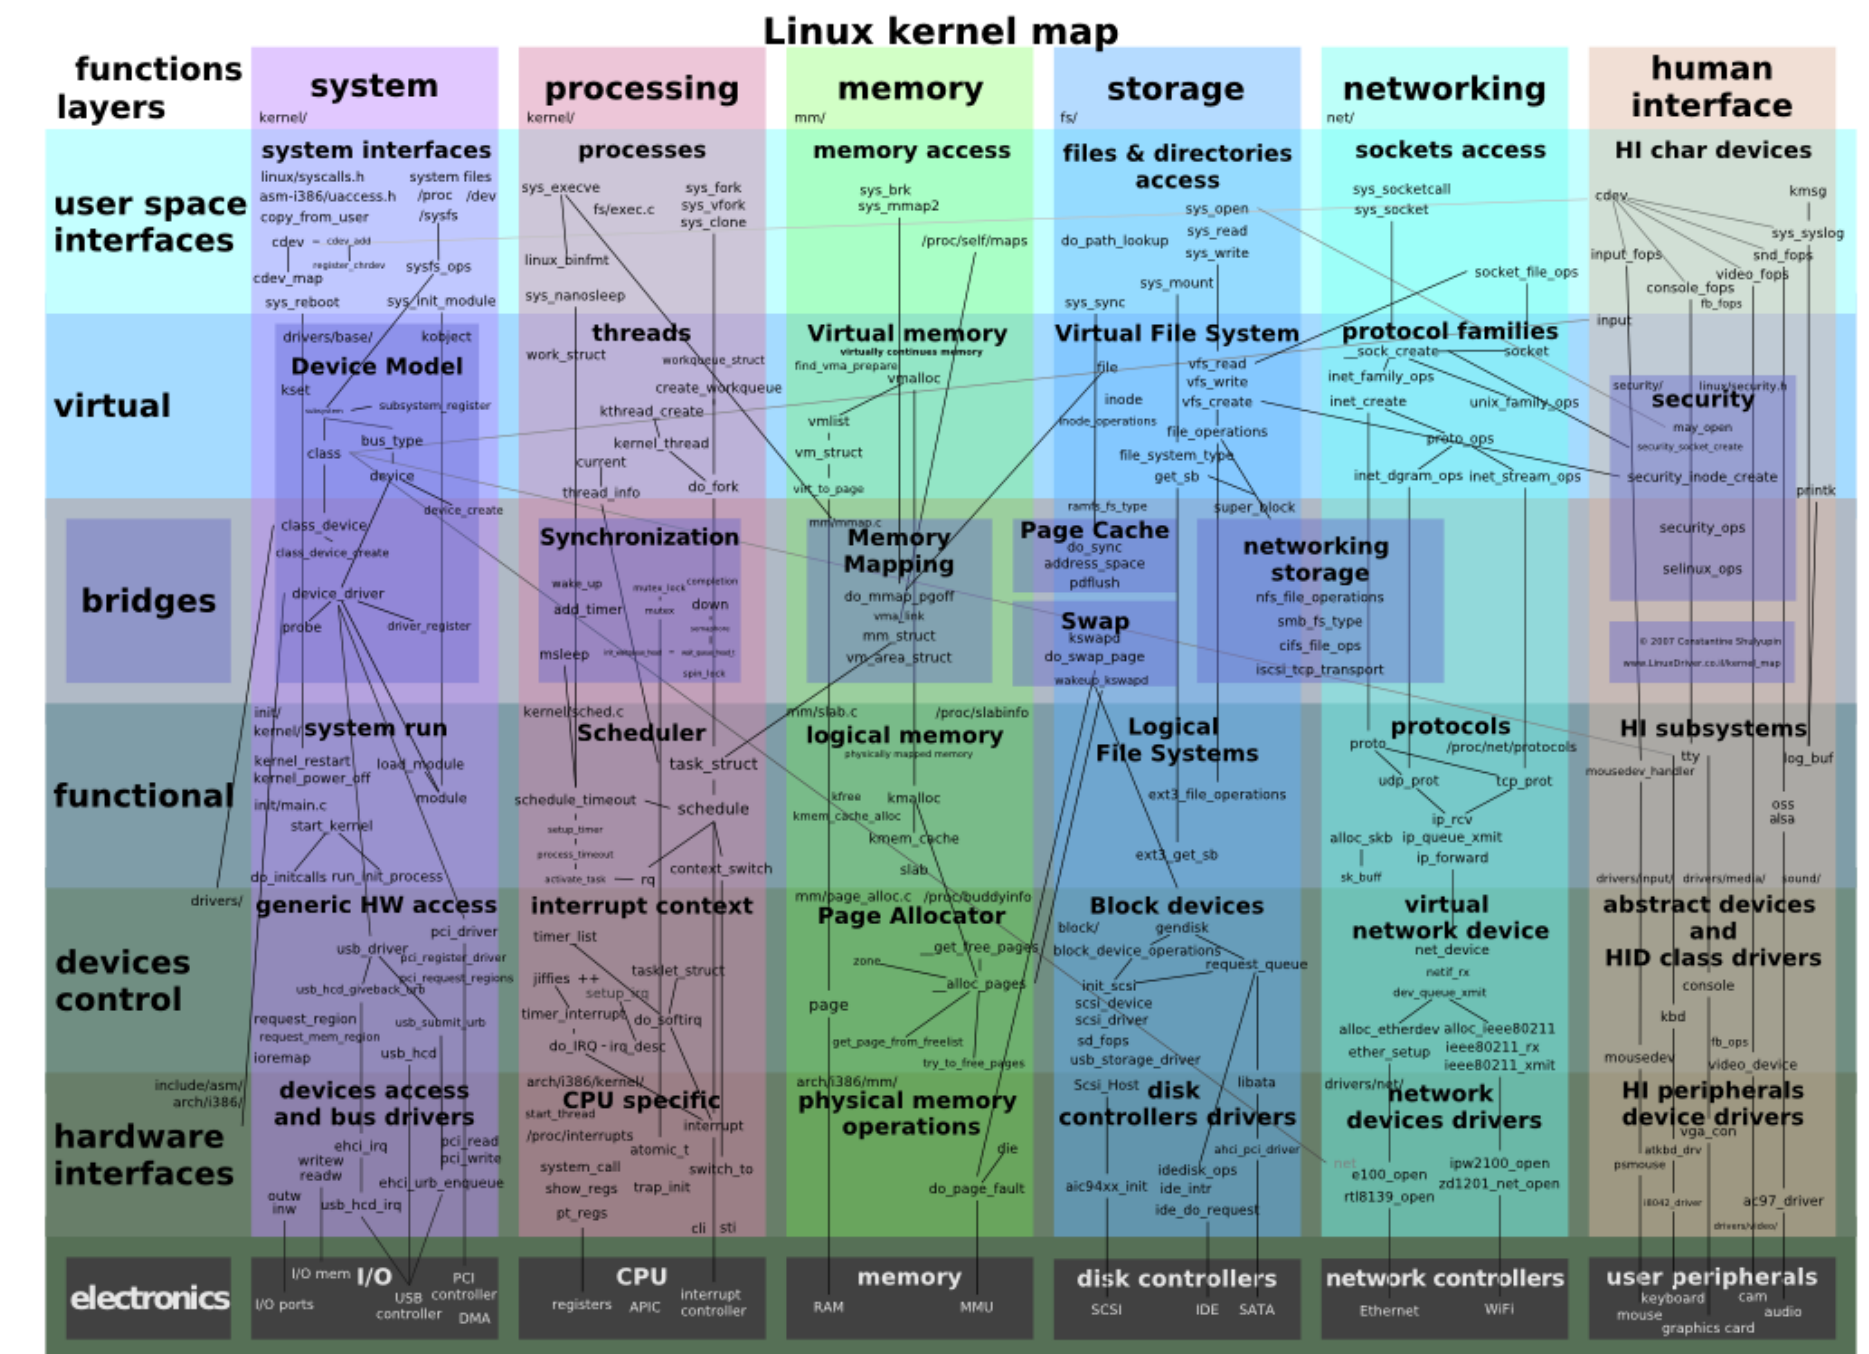
[Source](#)



[Source](#)

Initial Technology Choices (II)

- *Layer 2* (L2) -> machines joined by bridges.
- Bridge -> “slower” switch, “zeroconf”, part of the Linux kernel network stack.
- Virtual “wires” -> Virtual Ethernet Devices (veths), part of the kernel, manifest as regular NICs.
- Linux kernel -> great option:
 - Host machine -> L2 backbone.
 - L3 nodes -> “virtualised” network stack supporting veths.

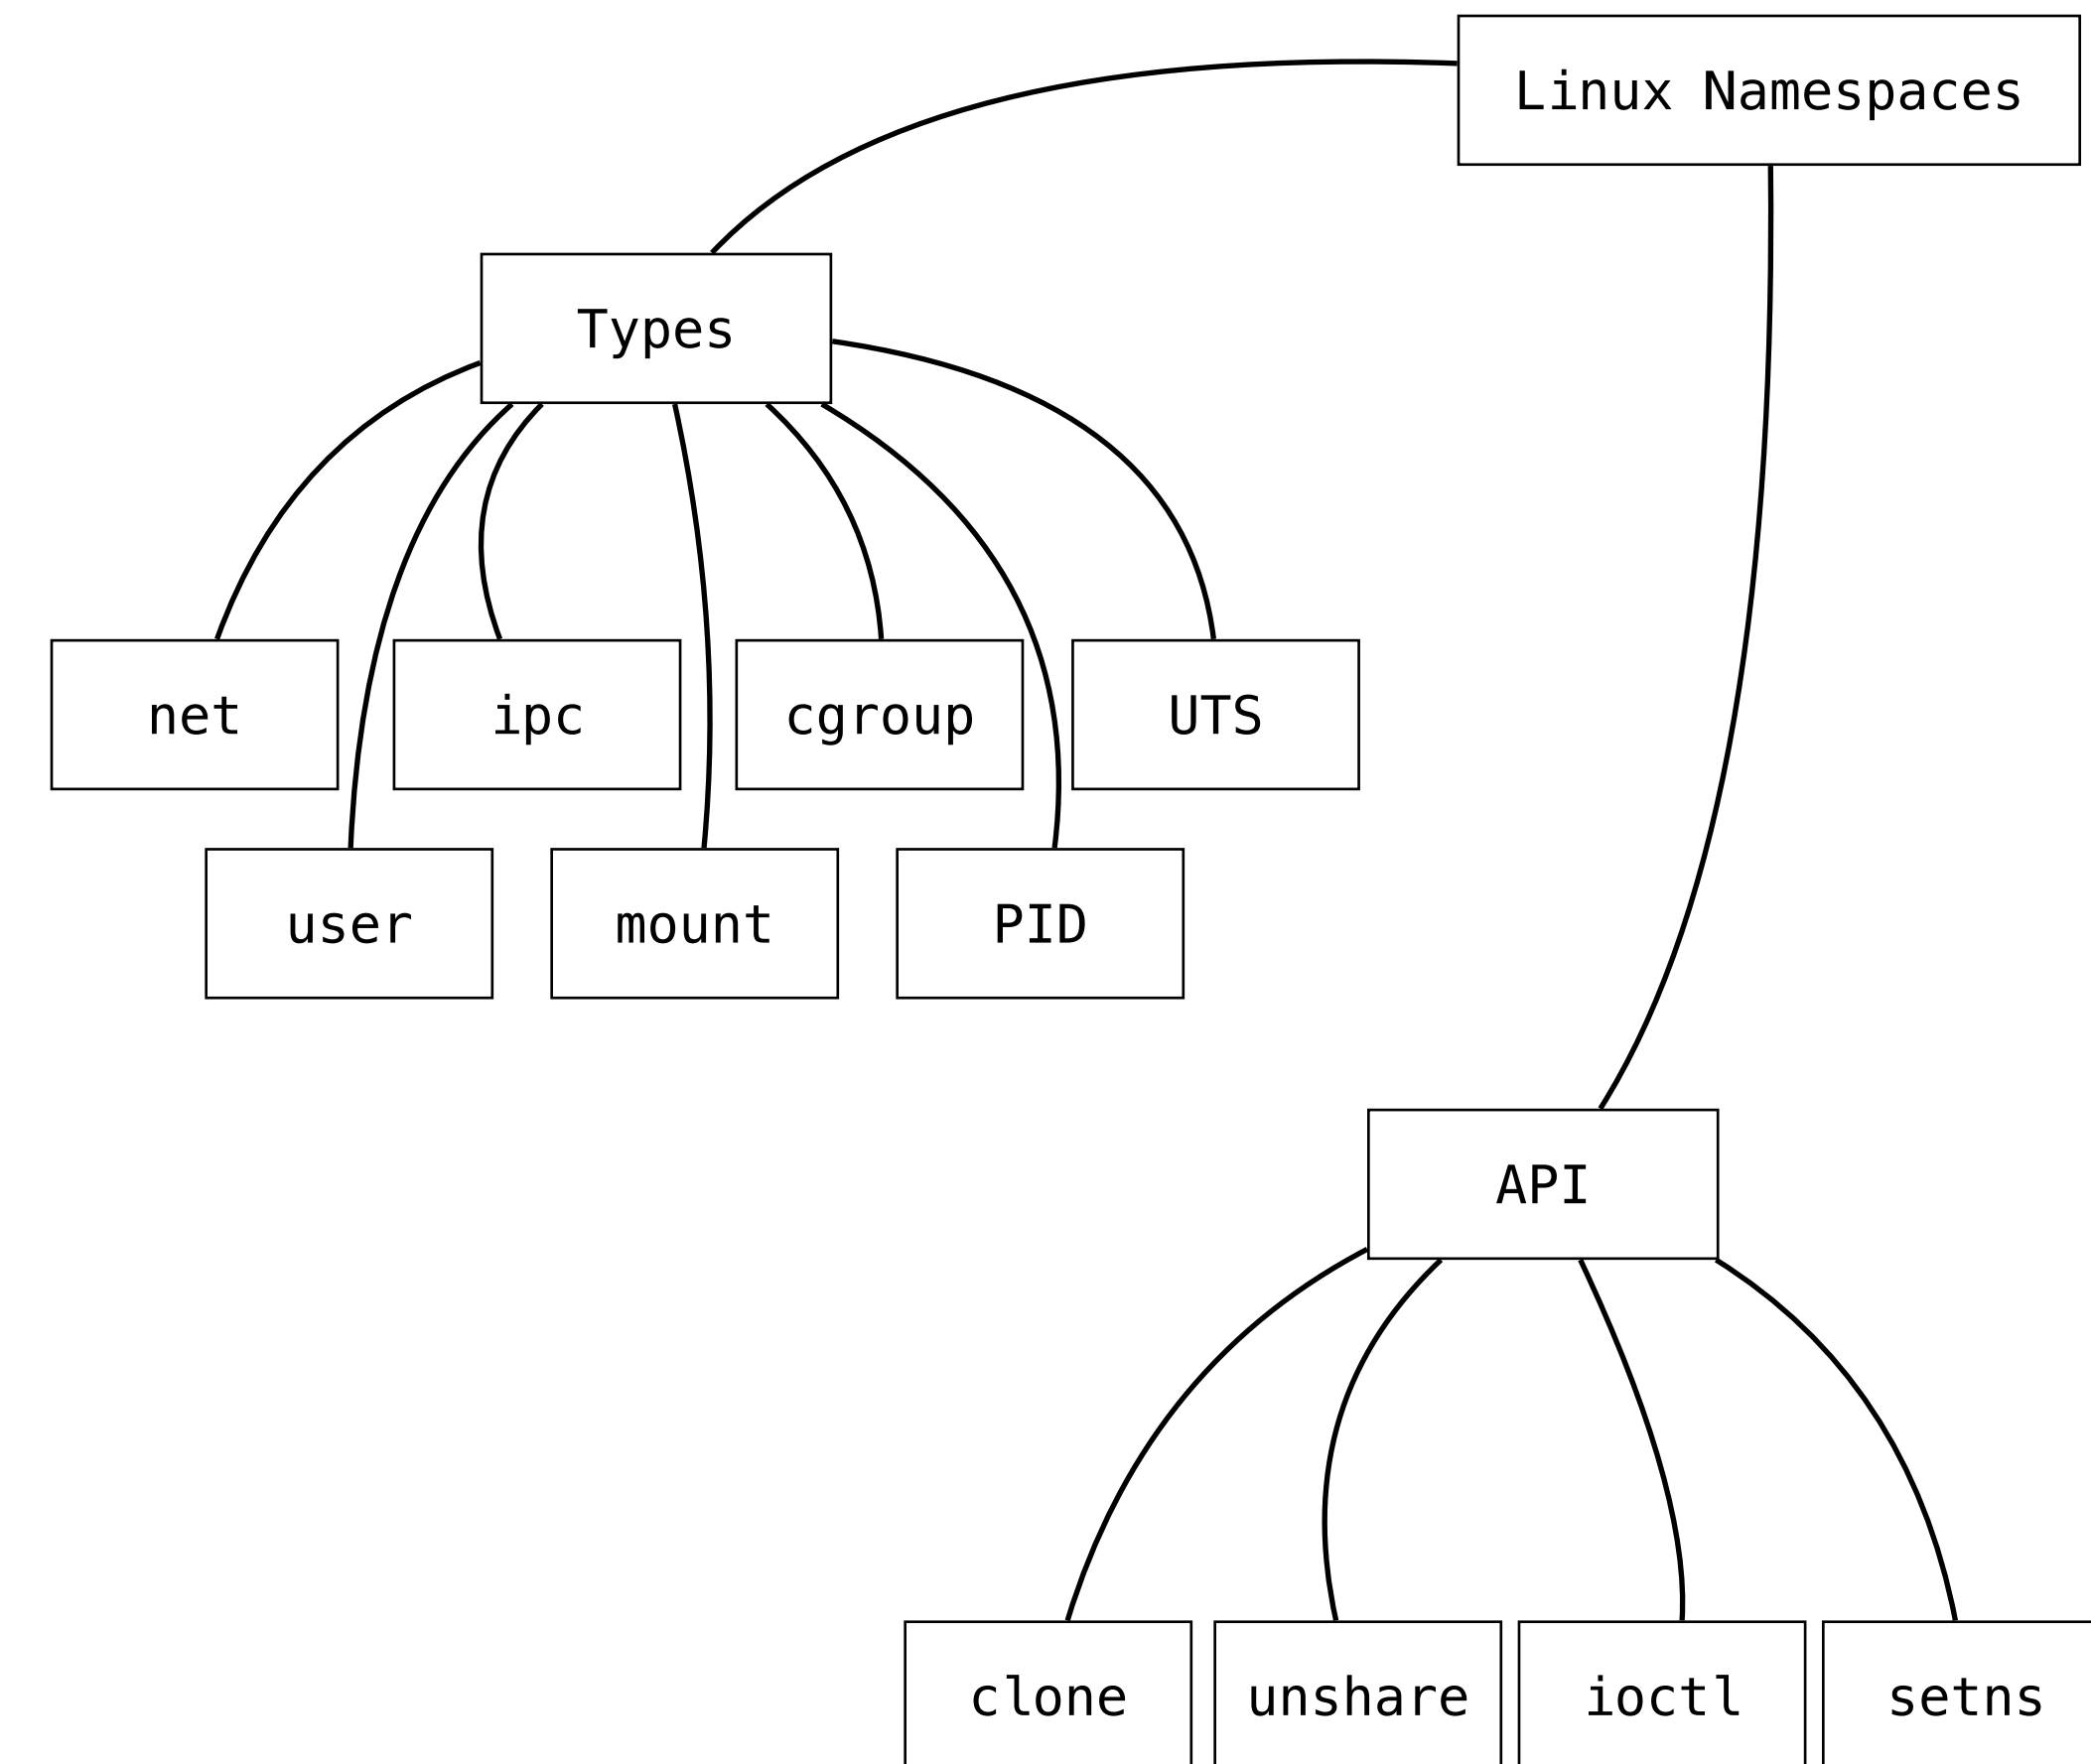


[Source](#)

- **namespaces** -> several network stacks on a single machine.

The Network Stack and the Namespace

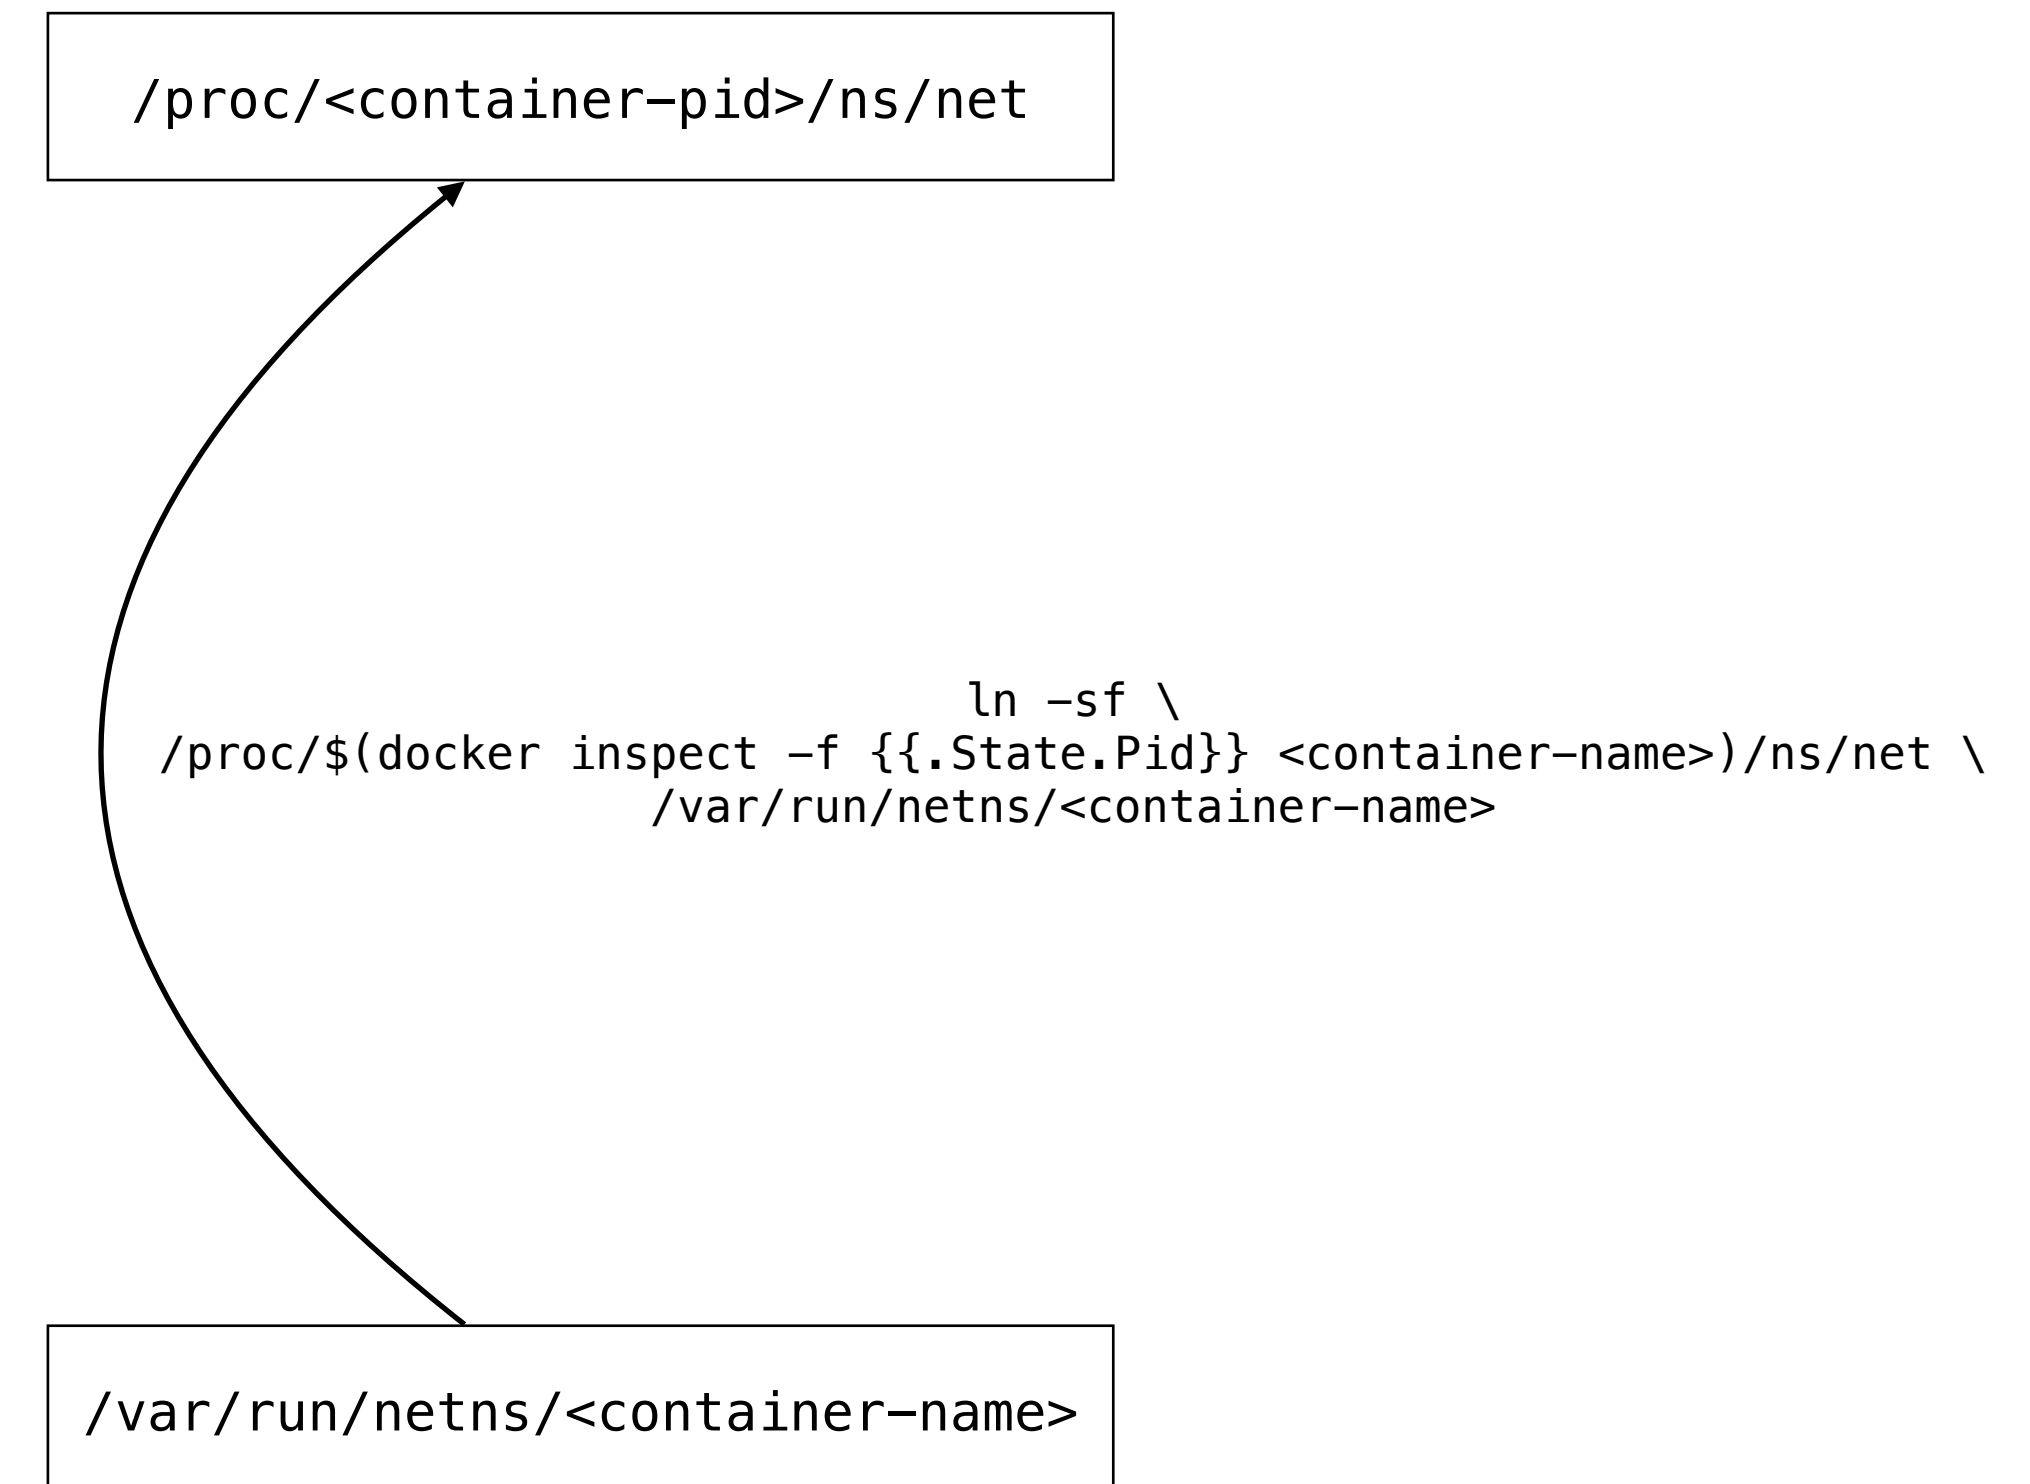
- Linux's network stack -> “blob” of code containing all the networking logic.
- Network stack -> logical entity in a network, usually 1 machine <-> 1 network stack.
- Virtual nodes -> independent network stack; two stacks = two virtual nodes.
- Namespaces -> partition kernel resources, openable objects.
- **Container -> its own network namespace.**



Based on [this image](#).

Interacting With the Network Stack

- iproute2 suite:
 - netlink interface with the kernel
 - Supersedes ifconfig
 - CLI interface with the user.
 - Instantiate and configure network -> on host and containers.
- iproute2 commands -> can be run from **any** namespace.
- Network namespaces **must be** linked under /var/run/netns/.



Instantiating Network Elements

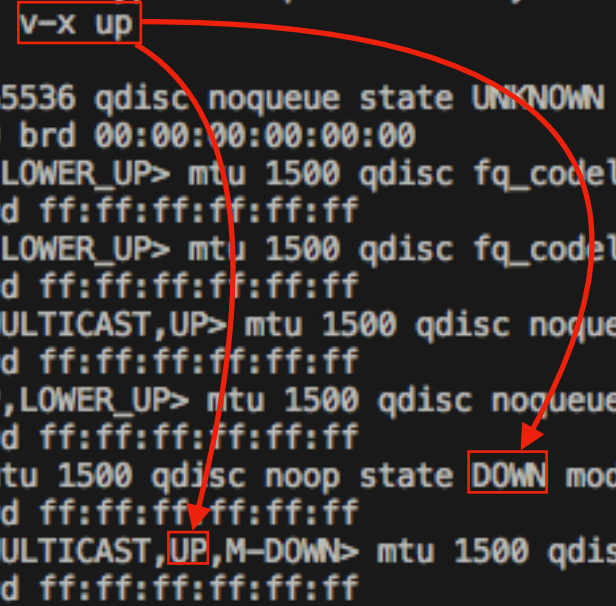
- **Bridges:**

- Creation: `ip link add foo-brd type bridge`
- Ignition: `ip link set foo-brd up`

- **Veths:**

- Creation: `ip link add v-x type veth peer name v-y`
- Ignition: `ip link set v-[x | y] up`
- Connection to a bridge: `ip link set v-[x | y] master foo-brd`
- Namespace change: `ip link set netns v-[x | y] netns-name`

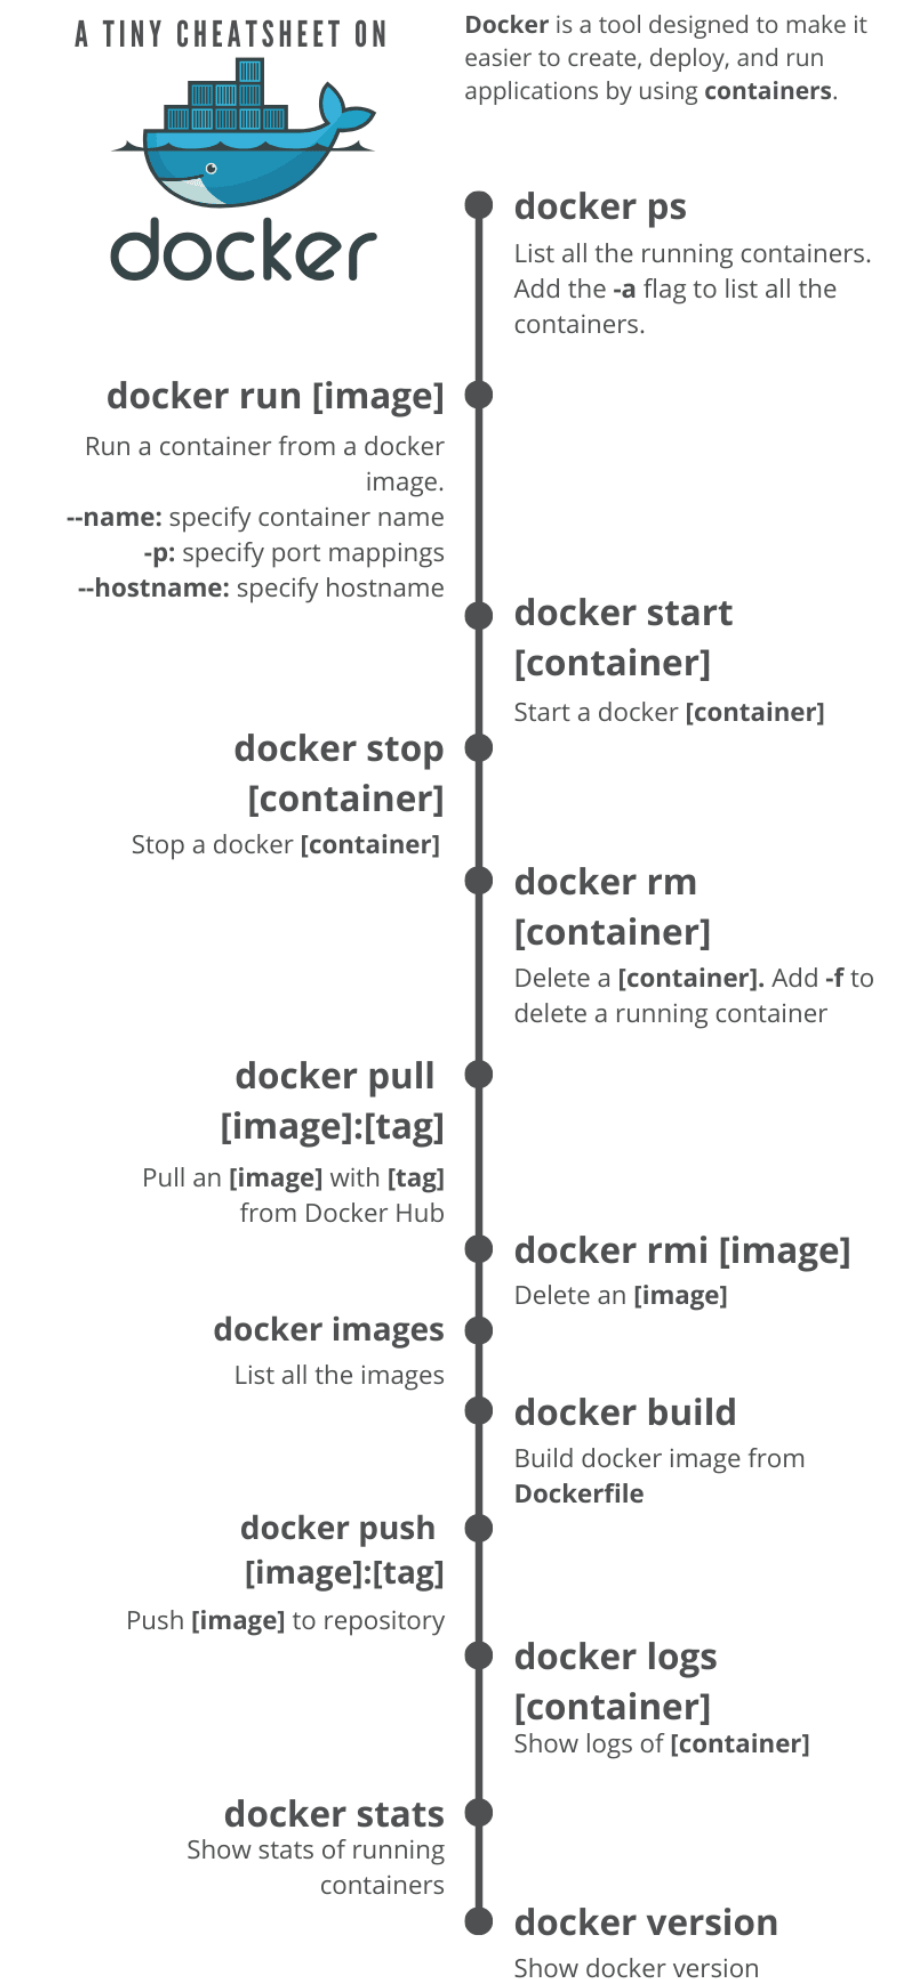
```
vagrant@focalvm:~$ sudo ip link add foo-brd type bridge
vagrant@focalvm:~$ sudo ip link set foo-brd up
vagrant@focalvm:~$ ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:b7:1d:9c:e0:75 brd ff:ff:ff:ff:ff:ff
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 08:00:27:3b:4f:dd brd ff:ff:ff:ff:ff:ff
4: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
    link/ether 02:42:6b:40:1e:4e brd ff:ff:ff:ff:ff:ff
7: foo-brd: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/ether c6:15:7f:ab:a6:ee brd ff:ff:ff:ff:ff:ff
vagrant@focalvm:~$ sudo ip link add v-x type veth peer name v-y
vagrant@focalvm:~$ sudo ip link set v-x up
vagrant@focalvm:~$ ip l
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:b7:1d:9c:e0:75 brd ff:ff:ff:ff:ff:ff
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 08:00:27:3b:4f:dd brd ff:ff:ff:ff:ff:ff
4: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
    link/ether 02:42:6b:40:1e:4e brd ff:ff:ff:ff:ff:ff
7: foo-brd: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/ether c6:15:7f:ab:a6:ee brd ff:ff:ff:ff:ff:ff
8: v-y@v-x: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether c6:3d:f9:fd:33:2f brd ff:ff:ff:ff:ff:ff
9: v-x@v-y: <NO-CARRIER,BROADCAST,MULTICAST,UP,M-DOWN> mtu 1500 qdisc noqueue state LOWERLAYERDOWN mode DEFAULT group default qlen 1000
    link/ether 06:af:a2:ff:26:32 brd ff:ff:ff:ff:ff:ff
vagrant@focalvm:~$
```



Instantiating Network Elements (II)

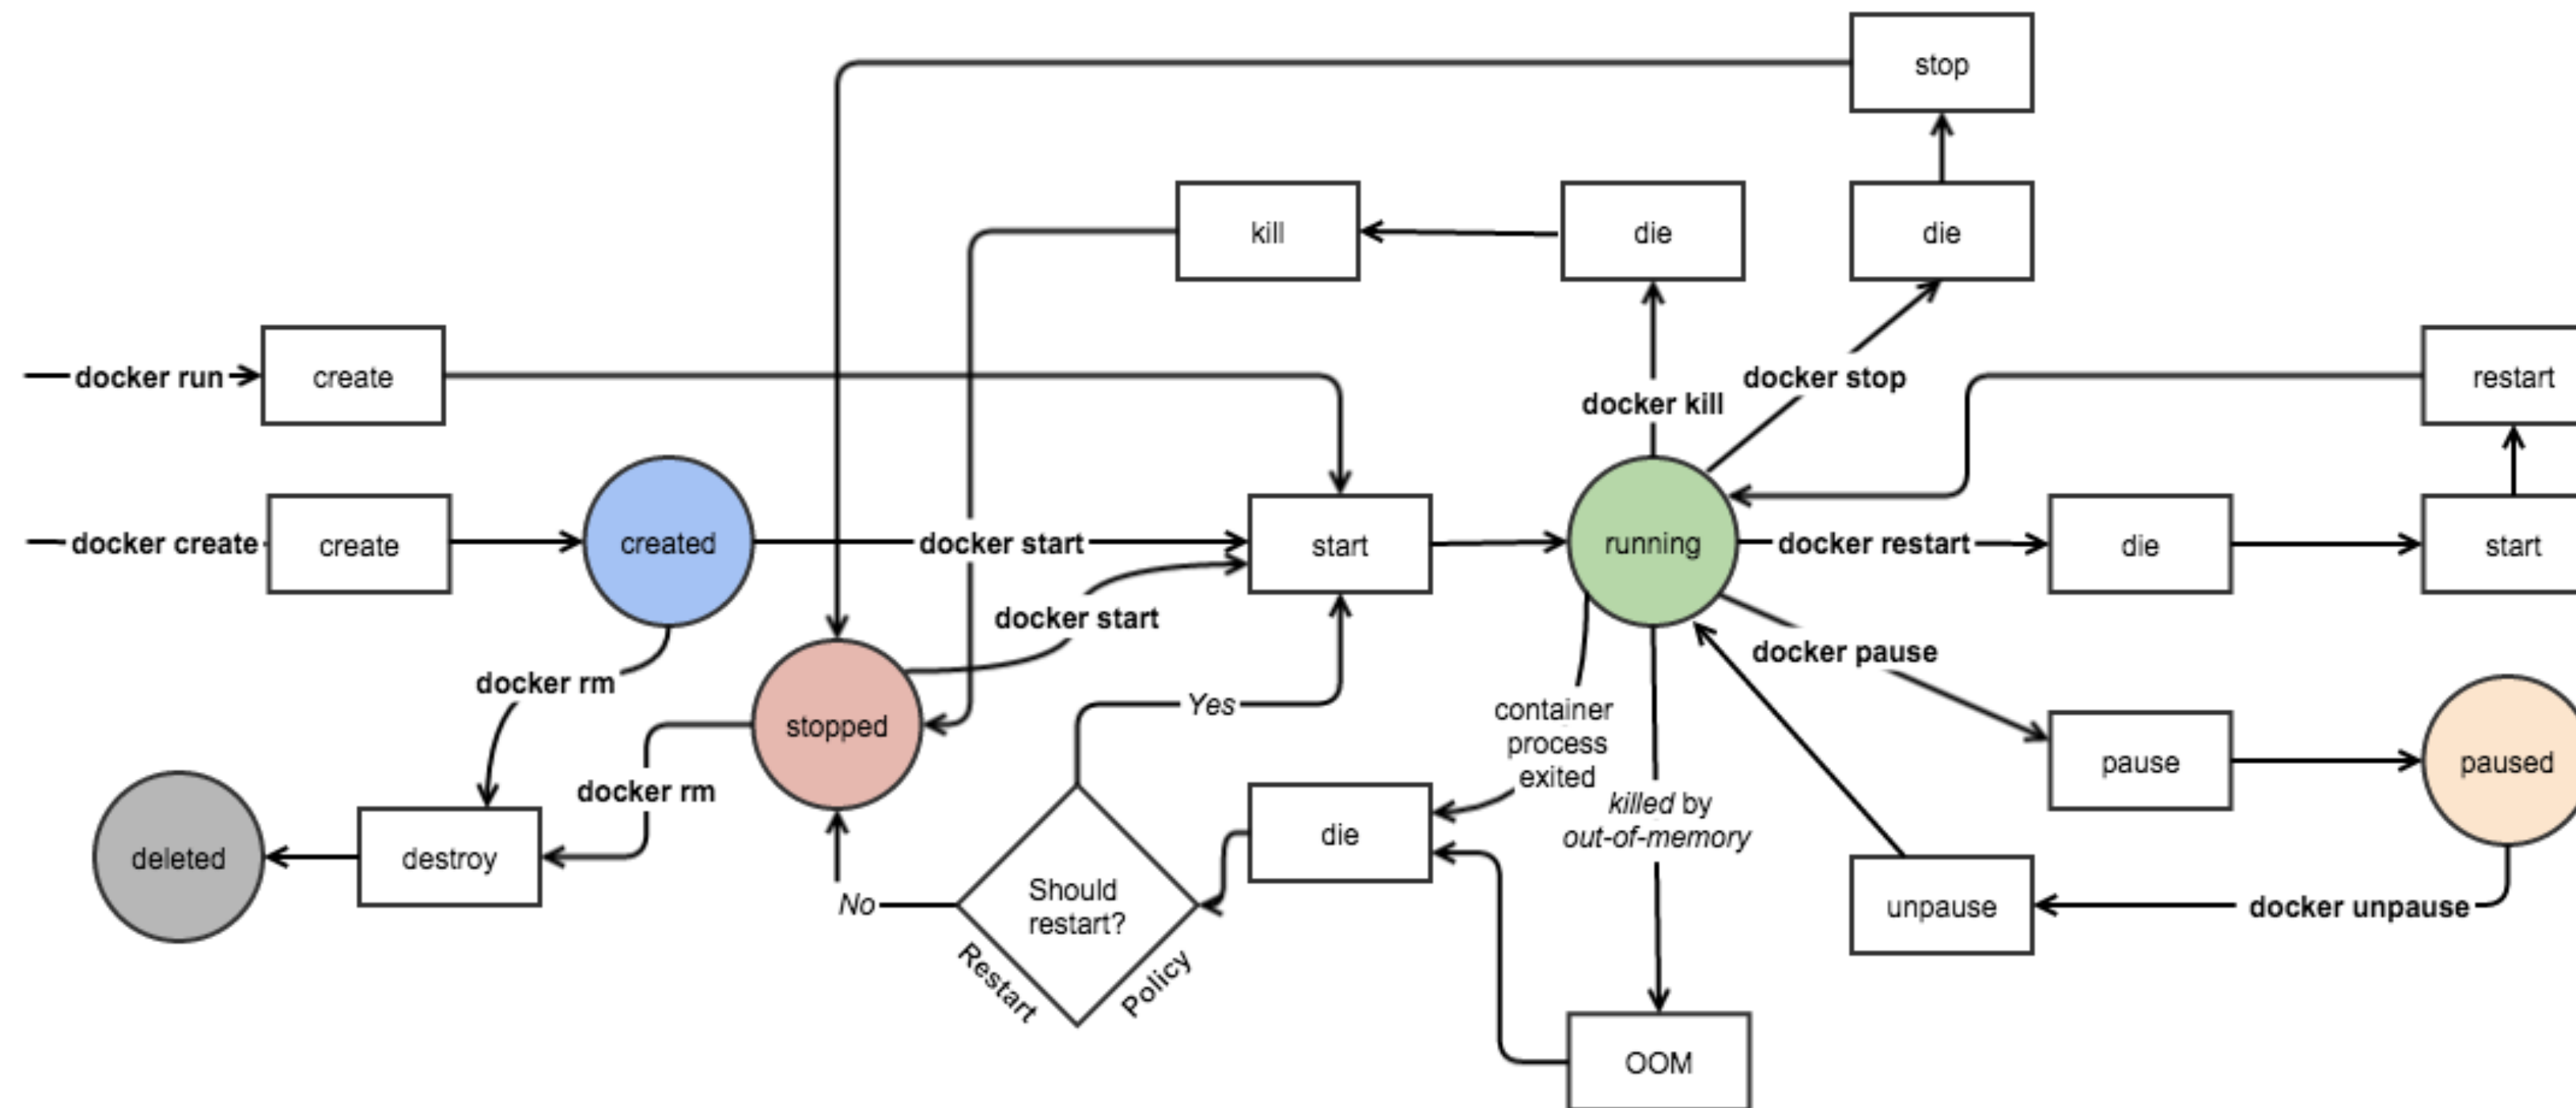
- Container -> runs an image.
- Dockerfiles
 - Define the contents of an image.
 - Are built to produce images.
- Built images -> run within a container.

```
docker run -d --name <container-name> --network none --cap-add ... <img-name>
```



A Container's Lifecycle

- `sshd` daemon -> keeps the container alive.
- `sshd` will not exit -> containers will run indefinitely.



Source

Routing and Addressing

- iproute2 -> manages addressing and routing.
- Option `-n <ns-name>`:
 - Run a command within any namespace.
 - Enables container configuration from the root namespace.
- Machines on same subnet -> automatic route.
- Address assignment:
 - `ip -n <cont> a add <IPv4>/<netmask> brd + dev <veth>`
- Route assignment:
 - `ip -n <cont> route add default via <gateway-IPv4>`
- Namespace -> associated container's name; configuration of containers through their name.

```
vagrant@focalvm:~$ docker exec -it r-z bash
root@r-z:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
44: veth_r-z_e-b@if43: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether f2:60:df:5b:4e:cc brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.2.1/24 brd 10.0.2.255 scope global veth_r-z_e-b
        valid_lft forever preferred_lft forever
57: veth_r-z_r-b@if56: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 72:49:24:30:ae:60 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.3.1/24 brd 10.0.3.255 scope global veth_r-z_r-b
        valid_lft forever preferred_lft forever
66: veth_r-z_v-b@if65: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether c2:0b:3d:26:cd:7b brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.4.1/24 brd 10.0.4.255 scope global veth_r-z_v-b
        valid_lft forever preferred_lft forever
75: veth_r-z_p-b@if74: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 3a:c2:0b:fb:c4:b9 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.5.1/24 brd 10.0.5.255 scope global veth_r-z_p-b
        valid_lft forever preferred_lft forever
90: veth_r-z_t-b-x@if89: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 1a:52:91:9c:ea:ae brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.6.1/24 brd 10.0.6.255 scope global veth_r-z_t-b-x
        valid_lft forever preferred_lft forever
110: veth_r-z_x-b@if109: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 2a:48:75:f4:bc:8d brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.8.2/24 brd 10.0.8.255 scope global veth_r-z_x-b
        valid_lft forever preferred_lft forever
root@r-z:/# ip r
10.0.0.0/24 via 10.0.8.1 dev veth_r-z_x-b
10.0.1.0/24 via 10.0.8.1 dev veth_r-z_x-b
10.0.2.0/24 dev veth_r-z_e-b proto kernel scope link src 10.0.2.1
10.0.3.0/24 dev veth_r-z_r-b proto kernel scope link src 10.0.3.1
10.0.4.0/24 dev veth_r-z_v-b proto kernel scope link src 10.0.4.1
10.0.5.0/24 dev veth_r-z_p-b proto kernel scope link src 10.0.5.1
10.0.6.0/24 dev veth_r-z_t-b-x proto kernel scope link src 10.0.6.1
10.0.8.0/24 dev veth_r-z_x-b proto kernel scope link src 10.0.8.2
root@r-z:/#
```


Firewall Configuration

- Topology needs -> forbid some connections.
- iptables -> chosen firewall implementation.
- Configuration -> based on chains defined by:
 - What packets are checked against it: INPUT, OUTPUT, FORWARD.
 - Default policy: DROP, ACCEPT.
 - Rules -> associated to a targets: DROP, ACCEPT.
- Firewall rules are applied within a given namespace.

```
vagrant@focalvm:~$ docker exec -it r-z iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source                destination

Chain FORWARD (policy DROP)
target     prot opt source                destination
ACCEPT    all  --  10.0.4.2                10.0.6.3
ACCEPT    all  --  10.0.6.3                10.0.4.2
ACCEPT    all  --  10.0.3.2                10.0.6.3
ACCEPT    all  --  10.0.6.3                10.0.3.2
ACCEPT    all  --  10.0.2.2                10.0.6.3
ACCEPT    all  --  10.0.6.3                10.0.2.2
ACCEPT    all  --  10.0.4.2                10.0.6.2
ACCEPT    all  --  10.0.6.2                10.0.4.2
ACCEPT    all  --  10.0.3.2                10.0.6.2
ACCEPT    all  --  10.0.6.2                10.0.3.2
ACCEPT    all  --  10.0.2.2                10.0.6.2
ACCEPT    all  --  10.0.6.2                10.0.2.2
ACCEPT    all  --  10.0.4.2                10.0.5.2
ACCEPT    all  --  10.0.5.2                10.0.4.2
ACCEPT    all  --  10.0.3.2                10.0.5.2
ACCEPT    all  --  10.0.5.2                10.0.3.2
ACCEPT    all  --  10.0.2.2                10.0.5.2
ACCEPT    all  --  10.0.5.2                10.0.2.2
ACCEPT    all  --  10.0.6.2                10.0.5.2
ACCEPT    all  --  10.0.5.2                10.0.6.2

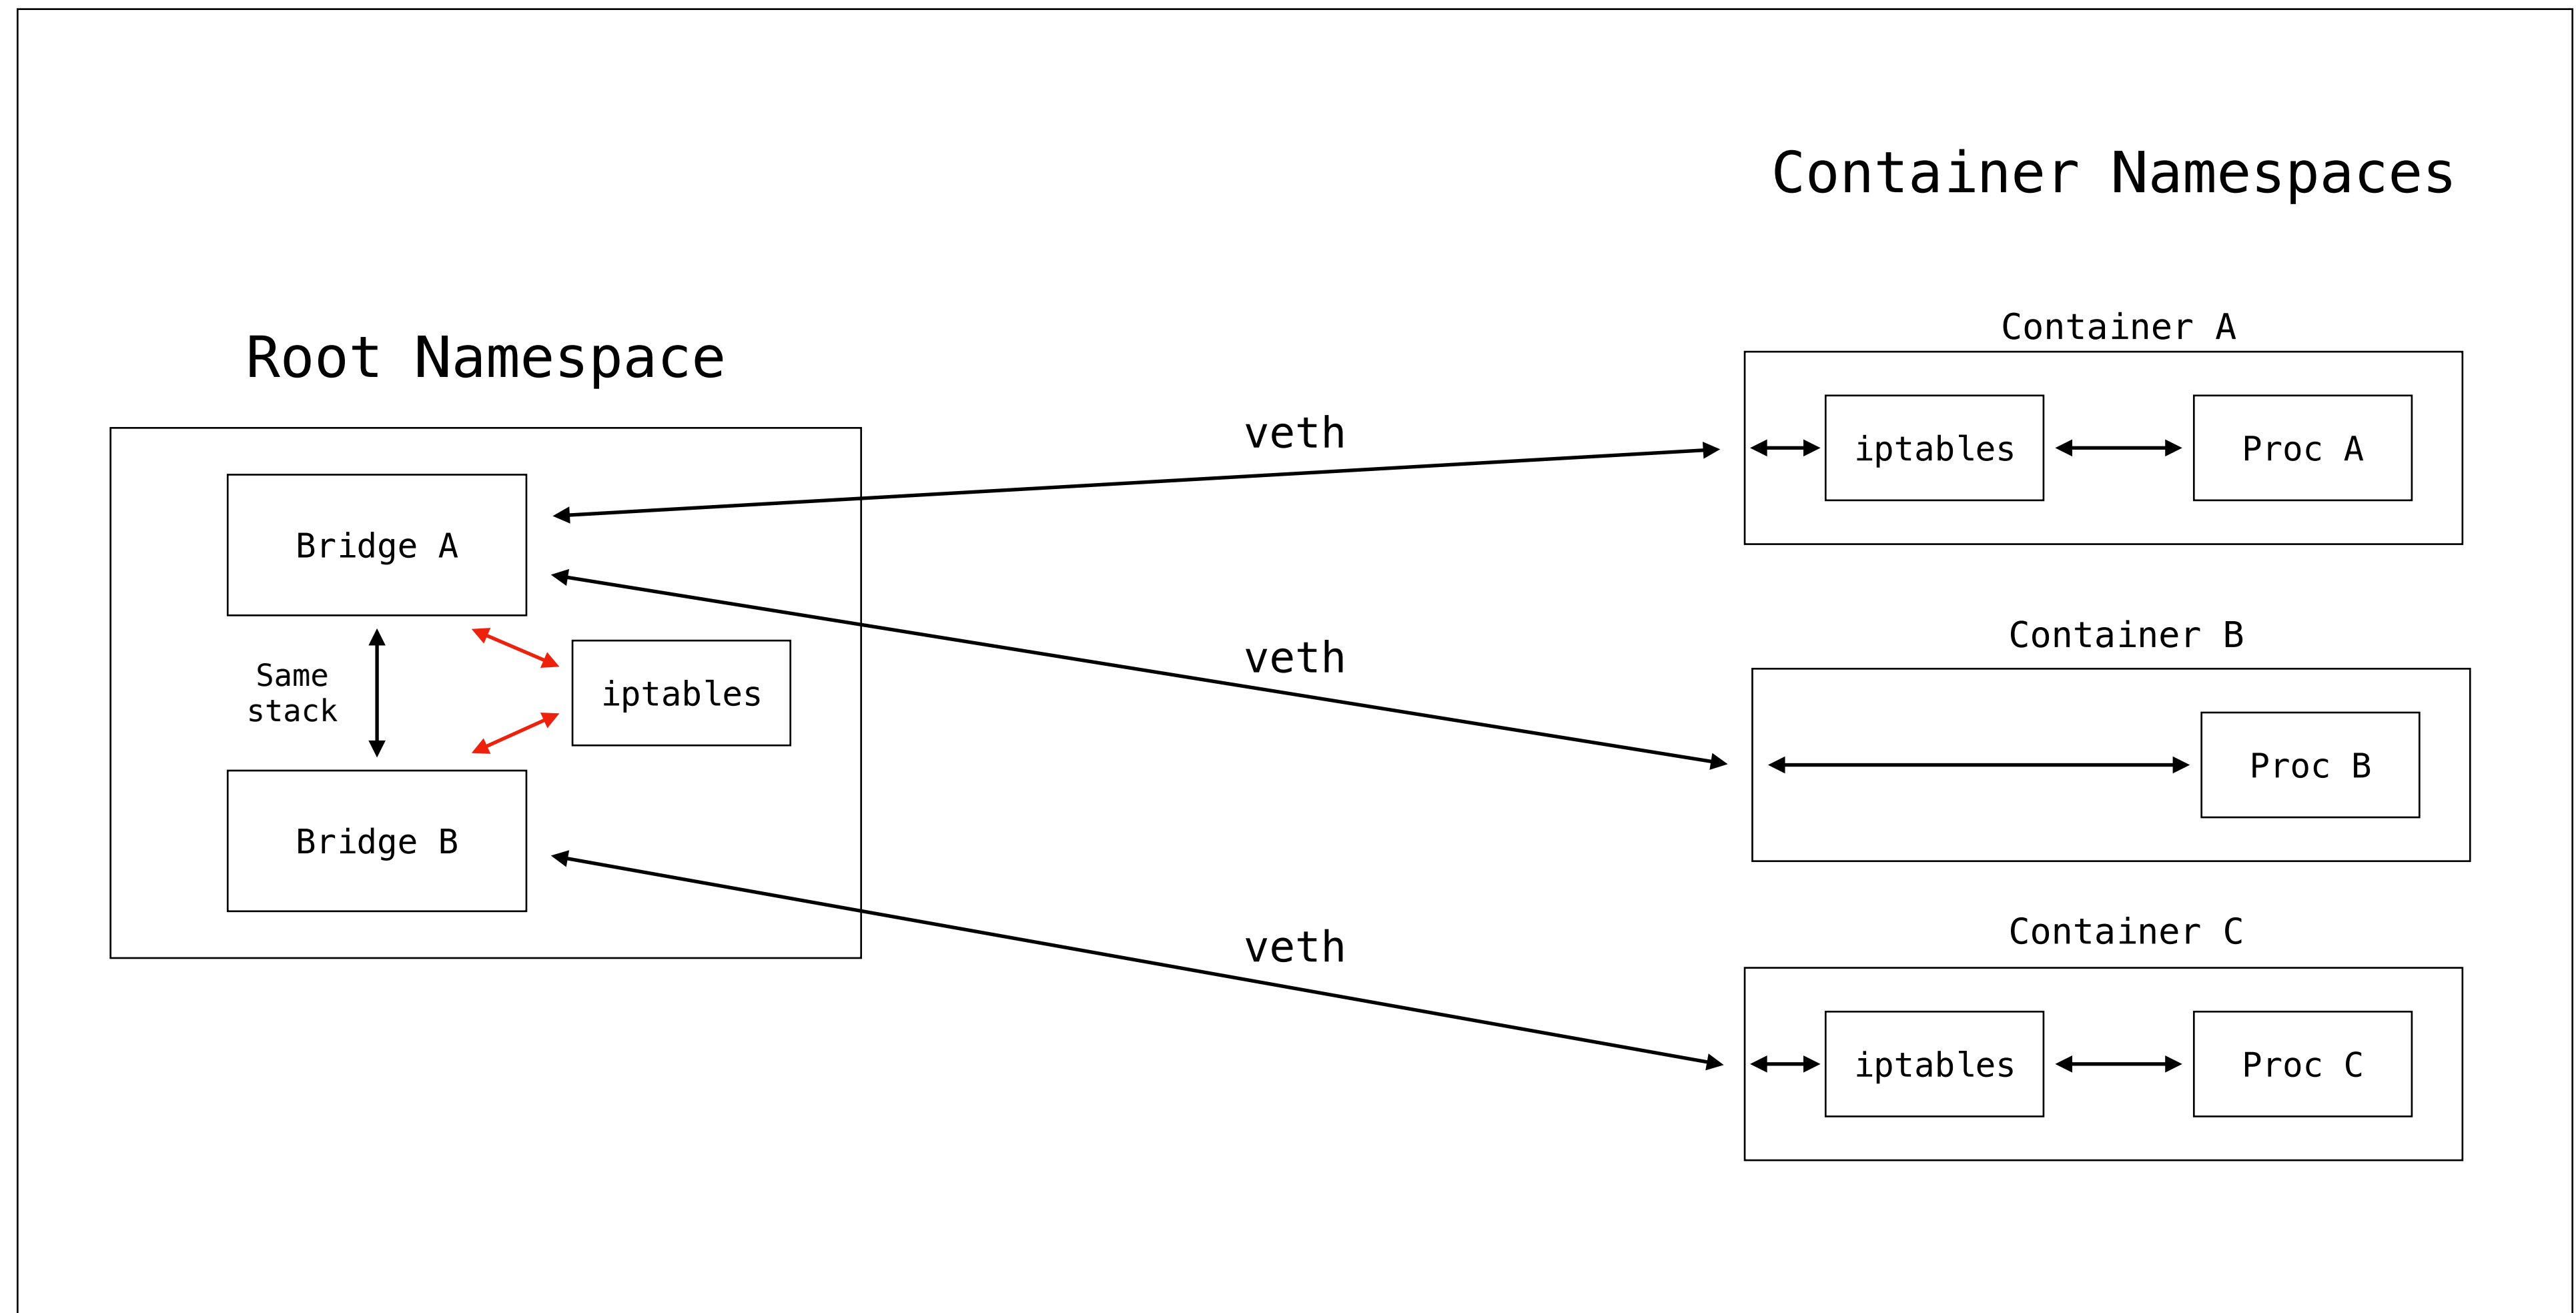
Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
vagrant@focalvm:~$
```

Subtleties

- Frames passing through bridges -> **checked** against iptables by default.
- Disabling the behaviour through sysctl:
 - Temporarily: `sysctl -w net.bridge.bridge-nf-call-iptables=0`
 - Permanently: Editing `/etc/sysctl.conf` and running `sysctl -p`.
- Machine -> **not allowed** to forward packets at L3 by default.
- Enabling the behaviour through sysctl:
 - Temporarily: `sysctl -w net.ipv4.ip_forward=1`
 - Permanently: Editing `/etc/sysctl.conf` and running `sysctl -p`.


“Physical” Topology

Physical Machine



A Simple Topology

- Listing 3.2 -> bash script instantiating a simple topology.
- Network's size increase -> script complexity and length grow quickly.
- Script -> no way of automatically altering the topology.
- System needs -> automatically instantiate and manage a network.



Sample Topology

- Subnet A -> 10.0.0.0/24
- Subnet B -> 10.0.1.0/24

H_{XY} -> Host Y belonging to subnet X


Our Tool

- Solution -> python3 program overcoming limitations.
- Leverages *Object Oriented Programming*.
- Interface for iproute2 and docker -> `os.system()`.
- NetworkX -> model networks as graphs: flexible network management.
- Alternative -> independent solution of NetworkX; less robust.
- User input -> graph representing desired network.
- User interface -> simple CLI.
- Basic commands:
 - mvnode: Allows to move a node to a new place within the network.
 - lsnet: Shows a graphical representation of the current network.

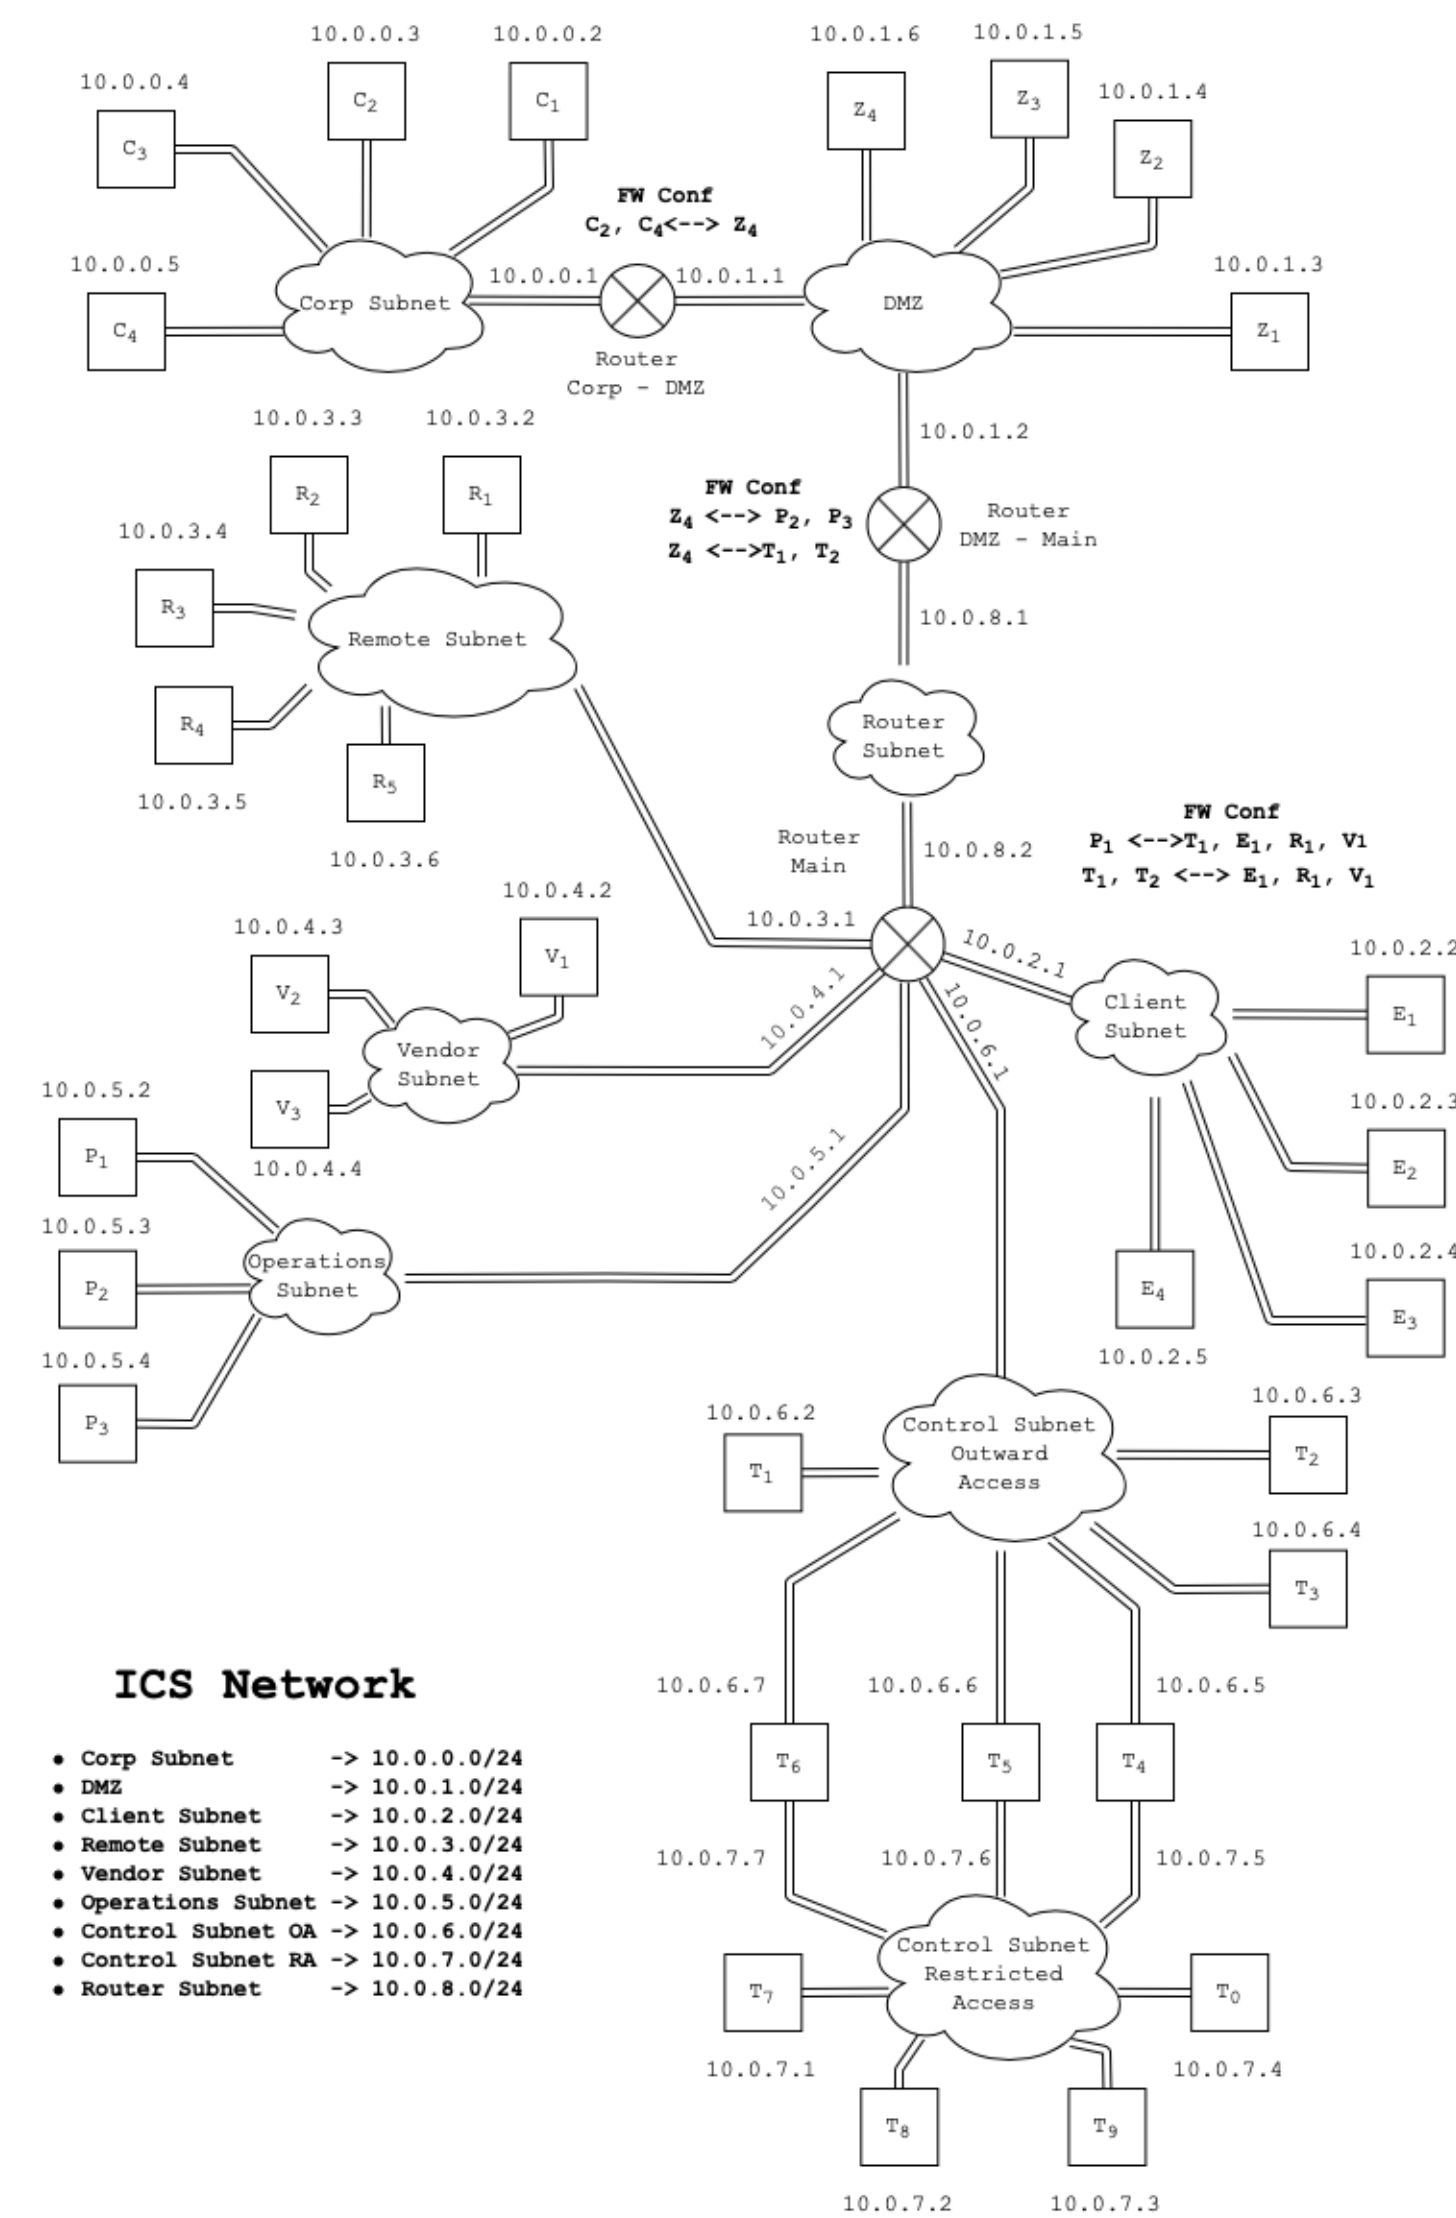
```
DBG: Veth interface veth_t-0_t-b-y is up!
DBG: Added veth pair veth_T-0_T-B-Y to bridge t-b-y
DBG: Veth interface veth_T-0_T-B-Y is up!
DBG: Bridge x-b is up!
DBG: Created veth pair veth_r-y_x-b-veth_R-Y_X-B
DBG: Added veth pair veth_r-y_x-b to node r-y
DBG: Veth interface veth_r-y_x-b is up!
DBG: Added veth pair veth_R-Y_X-B to bridge x-b
DBG: Veth interface veth_R-Y_X-B is up!
DBG: Created veth pair veth_r-z_x-b-veth_R-Z_X-B
DBG: Added veth pair veth_r-z_x-b to node r-z
DBG: Veth interface veth_r-z_x-b is up!
DBG: Added veth pair veth_R-Z_X-B to bridge x-b
DBG: Veth interface veth_R-Z_X-B is up!
DBG: Successfully assigned ip range 10.0.0.1/24 to interface veth_r-x_c-b belonging to host r-x
DBG: Successfully assigned ip range 10.0.1.1/24 to interface veth_r-x_z-b belonging to host r-x
DBG: Successfully assigned ip range 10.0.1.2/24 to interface veth_r-y_z-b belonging to host r-y
DBG: Successfully assigned ip range 10.0.8.1/24 to interface veth_r-y_x-b belonging to host r-y
DBG: Successfully assigned ip range 10.0.2.1/24 to interface veth_r-z_e-b belonging to host r-z
DBG: Successfully assigned ip range 10.0.3.1/24 to interface veth_r-z_r-b belonging to host r-z
DBG: Successfully assigned ip range 10.0.4.1/24 to interface veth_r-z_v-b belonging to host r-z
DBG: Successfully assigned ip range 10.0.5.1/24 to interface veth_r-z_p-b belonging to host r-z
DBG: Successfully assigned ip range 10.0.6.1/24 to interface veth_r-z_t-b-x belonging to host r-z
DBG: Successfully assigned ip range 10.0.8.2/24 to interface veth_r-z_x-b belonging to host r-z
DBG: Successfully assigned ip range 10.0.0.2/24 to interface veth_c-1_c-b belonging to host c-1
DBG: Successfully assigned ip range 10.0.0.3/24 to interface veth_c-2_c-b belonging to host c-2
DBG: Successfully assigned ip range 10.0.0.4/24 to interface veth_c-3_c-b belonging to host c-3
DBG: Successfully assigned ip range 10.0.0.5/24 to interface veth_c-4_c-b belonging to host c-4
DBG: Successfully assigned ip range 10.0.1.3/24 to interface veth_z-1_z-b belonging to host z-1
DBG: Successfully assigned ip range 10.0.1.4/24 to interface veth_z-2_z-b belonging to host z-2
DBG: Successfully assigned ip range 10.0.1.5/24 to interface veth_z-3_z-b belonging to host z-3
DBG: Successfully assigned ip range 10.0.1.6/24 to interface veth_z-4_z-b belonging to host z-4
DBG: Successfully assigned ip range 10.0.2.2/24 to interface veth_e-1_e-b belonging to host e-1
DBG: Successfully assigned ip range 10.0.2.3/24 to interface veth_e-2_e-b belonging to host e-2
DBG: Successfully assigned ip range 10.0.2.4/24 to interface veth_e-3_e-b belonging to host e-3
DBG: Successfully assigned ip range 10.0.2.5/24 to interface veth_e-4_e-b belonging to host e-4
DBG: Successfully assigned ip range 10.0.3.2/24 to interface veth_r-1_r-b belonging to host r-1
DBG: Successfully assigned ip range 10.0.3.3/24 to interface veth_r-2_r-b belonging to host r-2
DBG: Successfully assigned ip range 10.0.3.4/24 to interface veth_r-3_r-b belonging to host r-3
DBG: Successfully assigned ip range 10.0.3.5/24 to interface veth_r-4_r-b belonging to host r-4
DBG: Successfully assigned ip range 10.0.3.6/24 to interface veth_r-5_r-b belonging to host r-5
DBG: Successfully assigned ip range 10.0.4.2/24 to interface veth_v-1_v-b belonging to host v-1
DBG: Successfully assigned ip range 10.0.4.3/24 to interface veth_v-2_v-b belonging to host v-2
DBG: Successfully assigned ip range 10.0.4.4/24 to interface veth_v-3_v-b belonging to host v-3
DBG: Successfully assigned ip range 10.0.5.2/24 to interface veth_p-1_p-b belonging to host p-1
DBG: Successfully assigned ip range 10.0.5.3/24 to interface veth_p-2_p-b belonging to host p-2
DBG: Successfully assigned ip range 10.0.5.4/24 to interface veth_p-3_p-b belonging to host p-3
DBG: Successfully assigned ip range 10.0.6.2/24 to interface veth_t-1_t-b-x belonging to host t-1
DBG: Successfully assigned ip range 10.0.6.3/24 to interface veth_t-2_t-b-x belonging to host t-2
DBG: Successfully assigned ip range 10.0.6.4/24 to interface veth_t-3_t-b-x belonging to host t-3
DBG: Successfully assigned ip range 10.0.6.5/24 to interface veth_t-4_t-b-x belonging to host t-4
DBG: Successfully assigned ip range 10.0.6.6/24 to interface veth_t-5_t-b-x belonging to host t-5
DBG: Successfully assigned ip range 10.0.6.7/24 to interface veth_t-6_t-b-x belonging to host t-6
DBG: Successfully assigned ip range 10.0.7.1/24 to interface veth_t-4_t-b-y belonging to host t-4
DBG: Successfully assigned ip range 10.0.7.2/24 to interface veth_t-5_t-b-y belonging to host t-5
DBG: Successfully assigned ip range 10.0.7.3/24 to interface veth_t-6_t-b-y belonging to host t-6
DBG: Successfully assigned ip range 10.0.7.4/24 to interface veth_t-7_t-b-y belonging to host t-7
DBG: Successfully assigned ip range 10.0.7.5/24 to interface veth_t-8_t-b-y belonging to host t-8
DBG: Successfully assigned ip range 10.0.7.6/24 to interface veth_t-9_t-b-y belonging to host t-9
DBG: Successfully assigned ip range 10.0.7.7/24 to interface veth_t-0_t-b-y belonging to host t-0
DBG: Successfully added route to 10.0.0.0/24 through 10.0.1.1 at z-1
DBG: Successfully added route to 10.0.0.0/24 through 10.0.1.1 at z-2
DBG: Successfully added route to 10.0.0.0/24 through 10.0.1.1 at z-3
DBG: Successfully added route to 10.0.0.0/24 through 10.0.1.1 at z-4
DBG: Successfully added route to 10.0.0.0/24 through 10.0.1.1 at r-y
DBG: Successfully added route to 10.0.0.0/24 through 10.0.8.1 at r-z
DBG: Successfully added route to 10.0.0.0/24 through 10.0.3.1 at r-1
DBG: Successfully added route to 10.0.0.0/24 through 10.0.3.1 at r-2
DBG: Successfully added route to 10.0.0.0/24 through 10.0.3.1 at r-3
```


Tool Modules

- Tool architecture -> 3 separate modules:
 - **virt_net**
 - Instantiation, configuration and addressing of network elements.
 - Classes representing network elements: veths and routers.
 - **graph_interpreter**
 - Translates graphs provided by users to classes defined in virt_net.
 - Contains the logic of features such as:
 - Node movement.
 - Network's routing.
 - **net_ctrl**
 - Implements the CLI presented to users.
 - Commands are converted to calls to the graph_interpreter module.



Our Most Complex Topology



Based on figure 4 of “Improving ICS cyber resilience through optimal diversification of network resources”. [Online article](#).

Proof of Concept

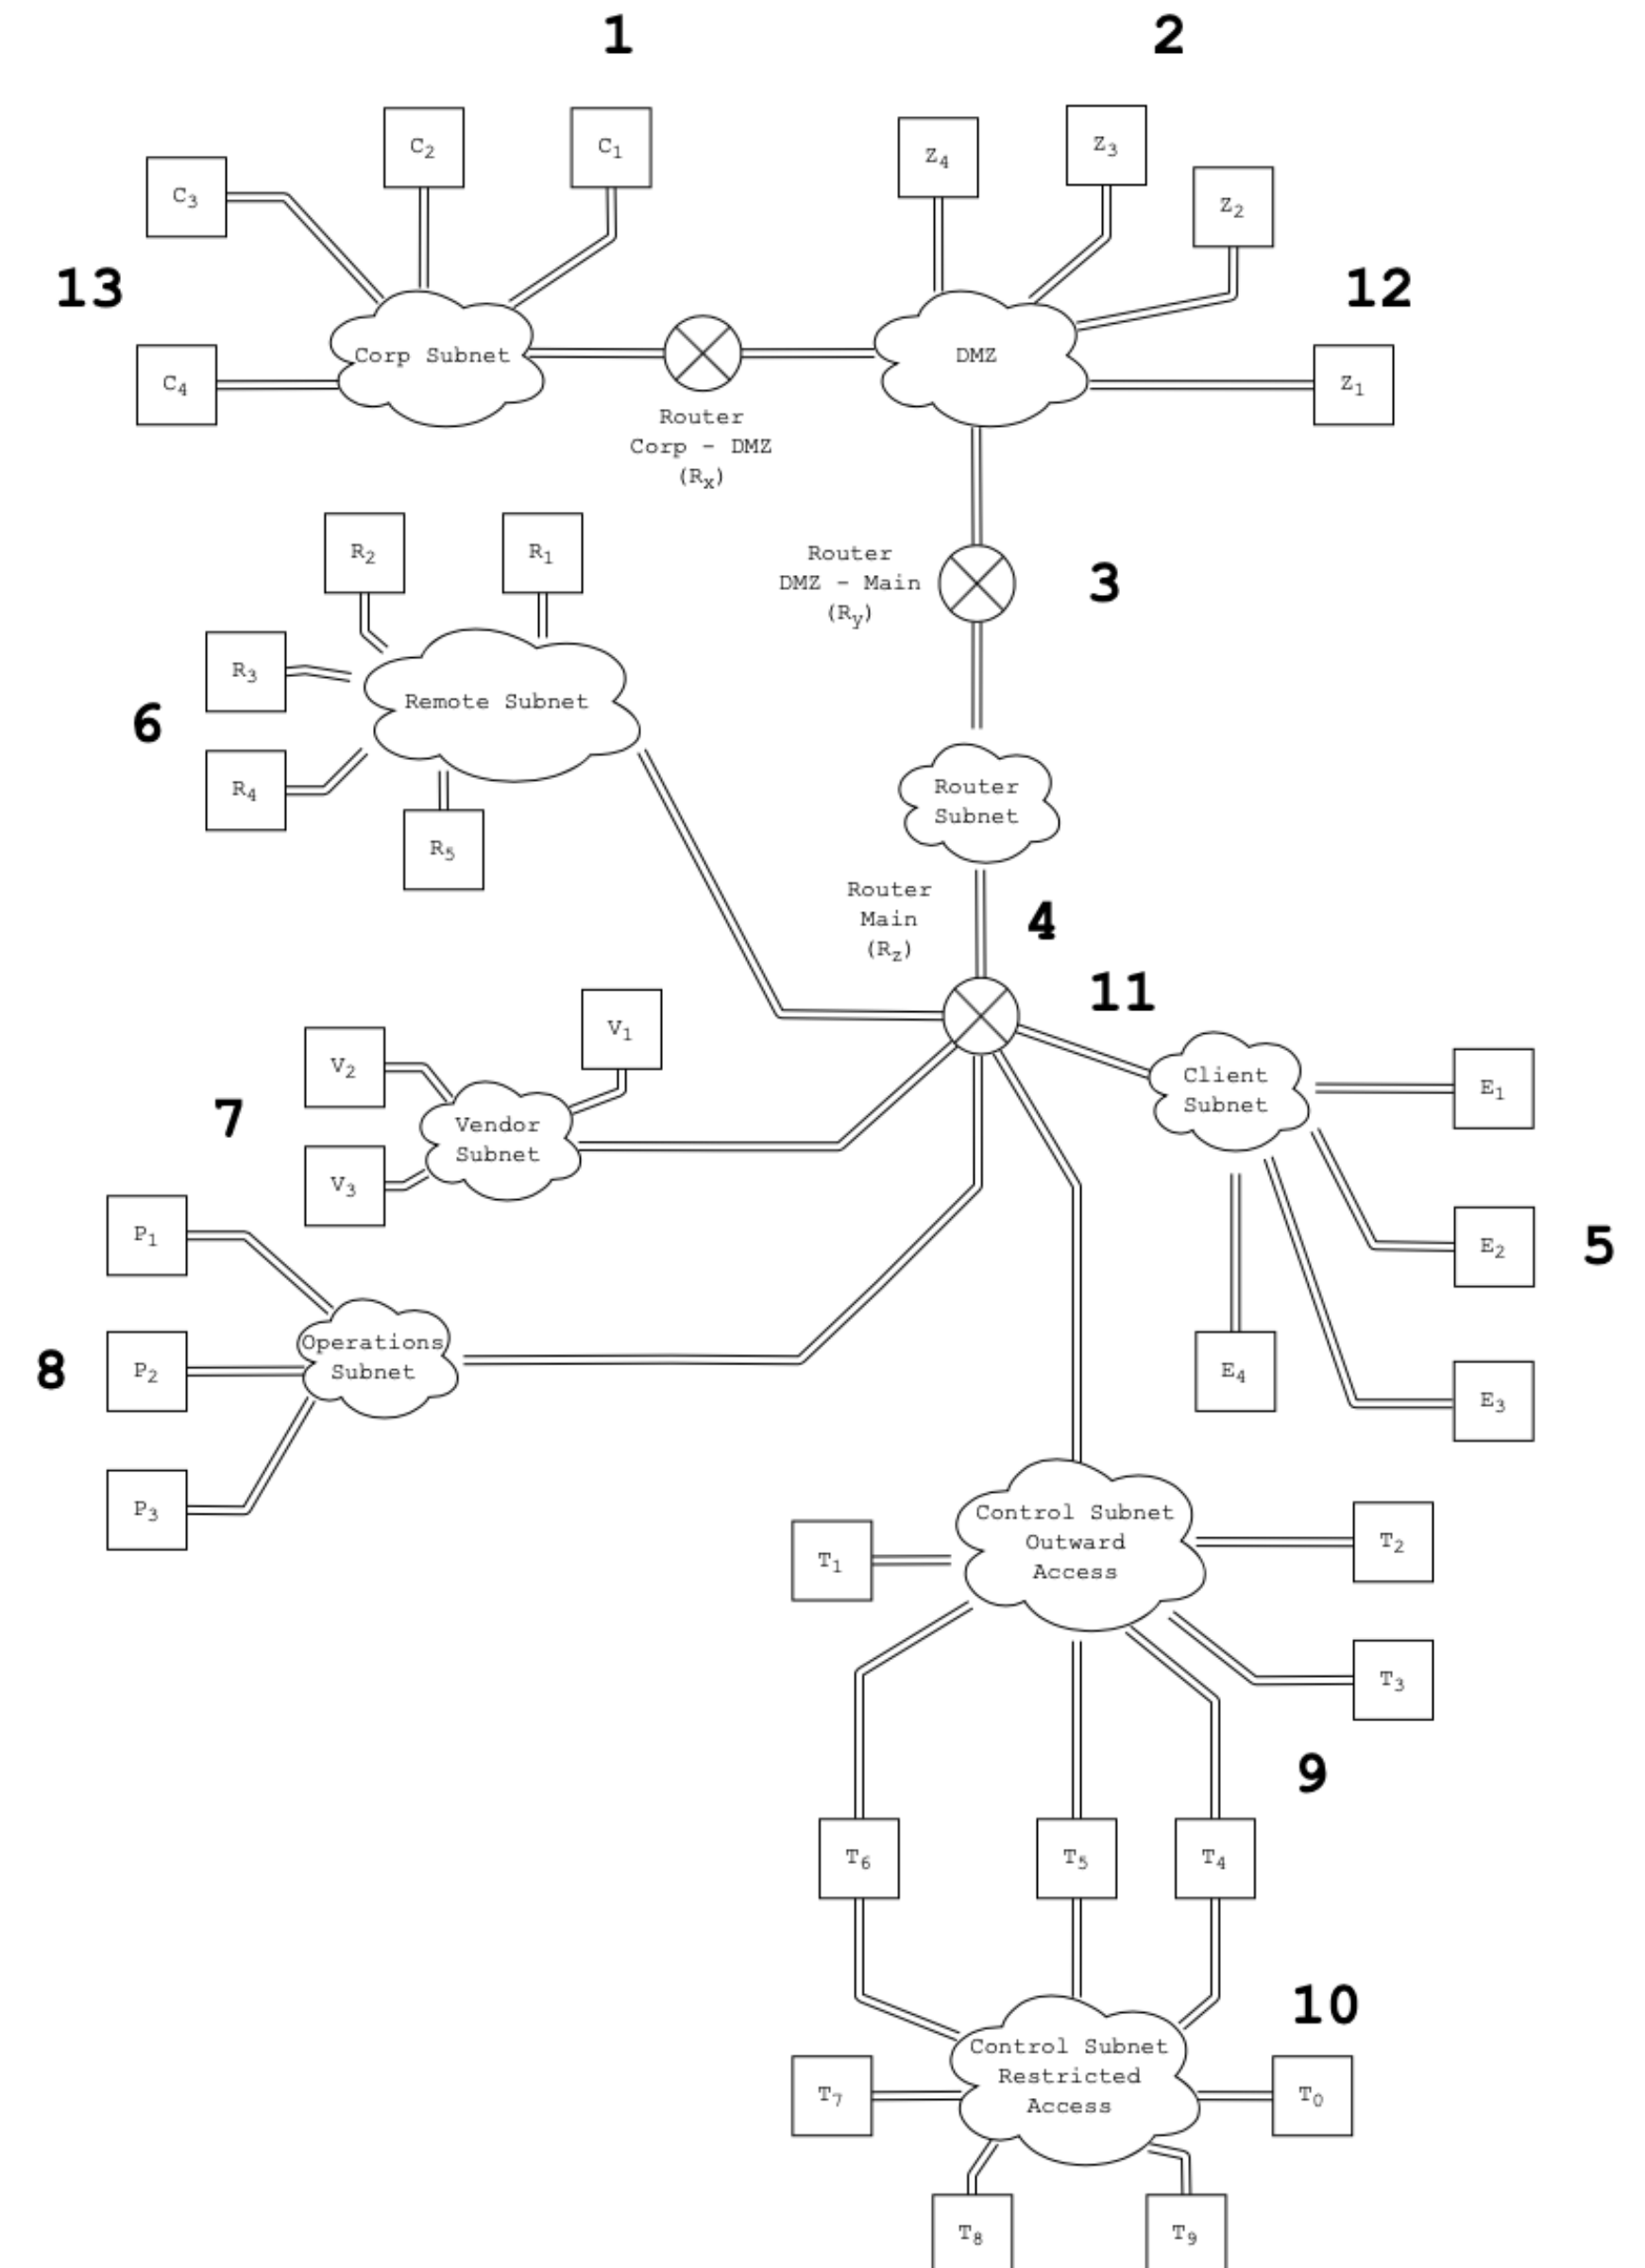
- Simulate an attack -> monitor the network's state during and after.
- Objective -> show our tool can be used for its intended purpose.
- We define the *Quality of Service* of the network as:

$$QoS(t) = \frac{pings(t)}{pings(0)}; \quad QoS(t) \in [0, 1] \quad \forall t \in [0, \infty)$$


- Attack -> diminishes the QoS.
- Information of interest -> effect of topology changes on the QoS evolution.

Proof of Concept (II)

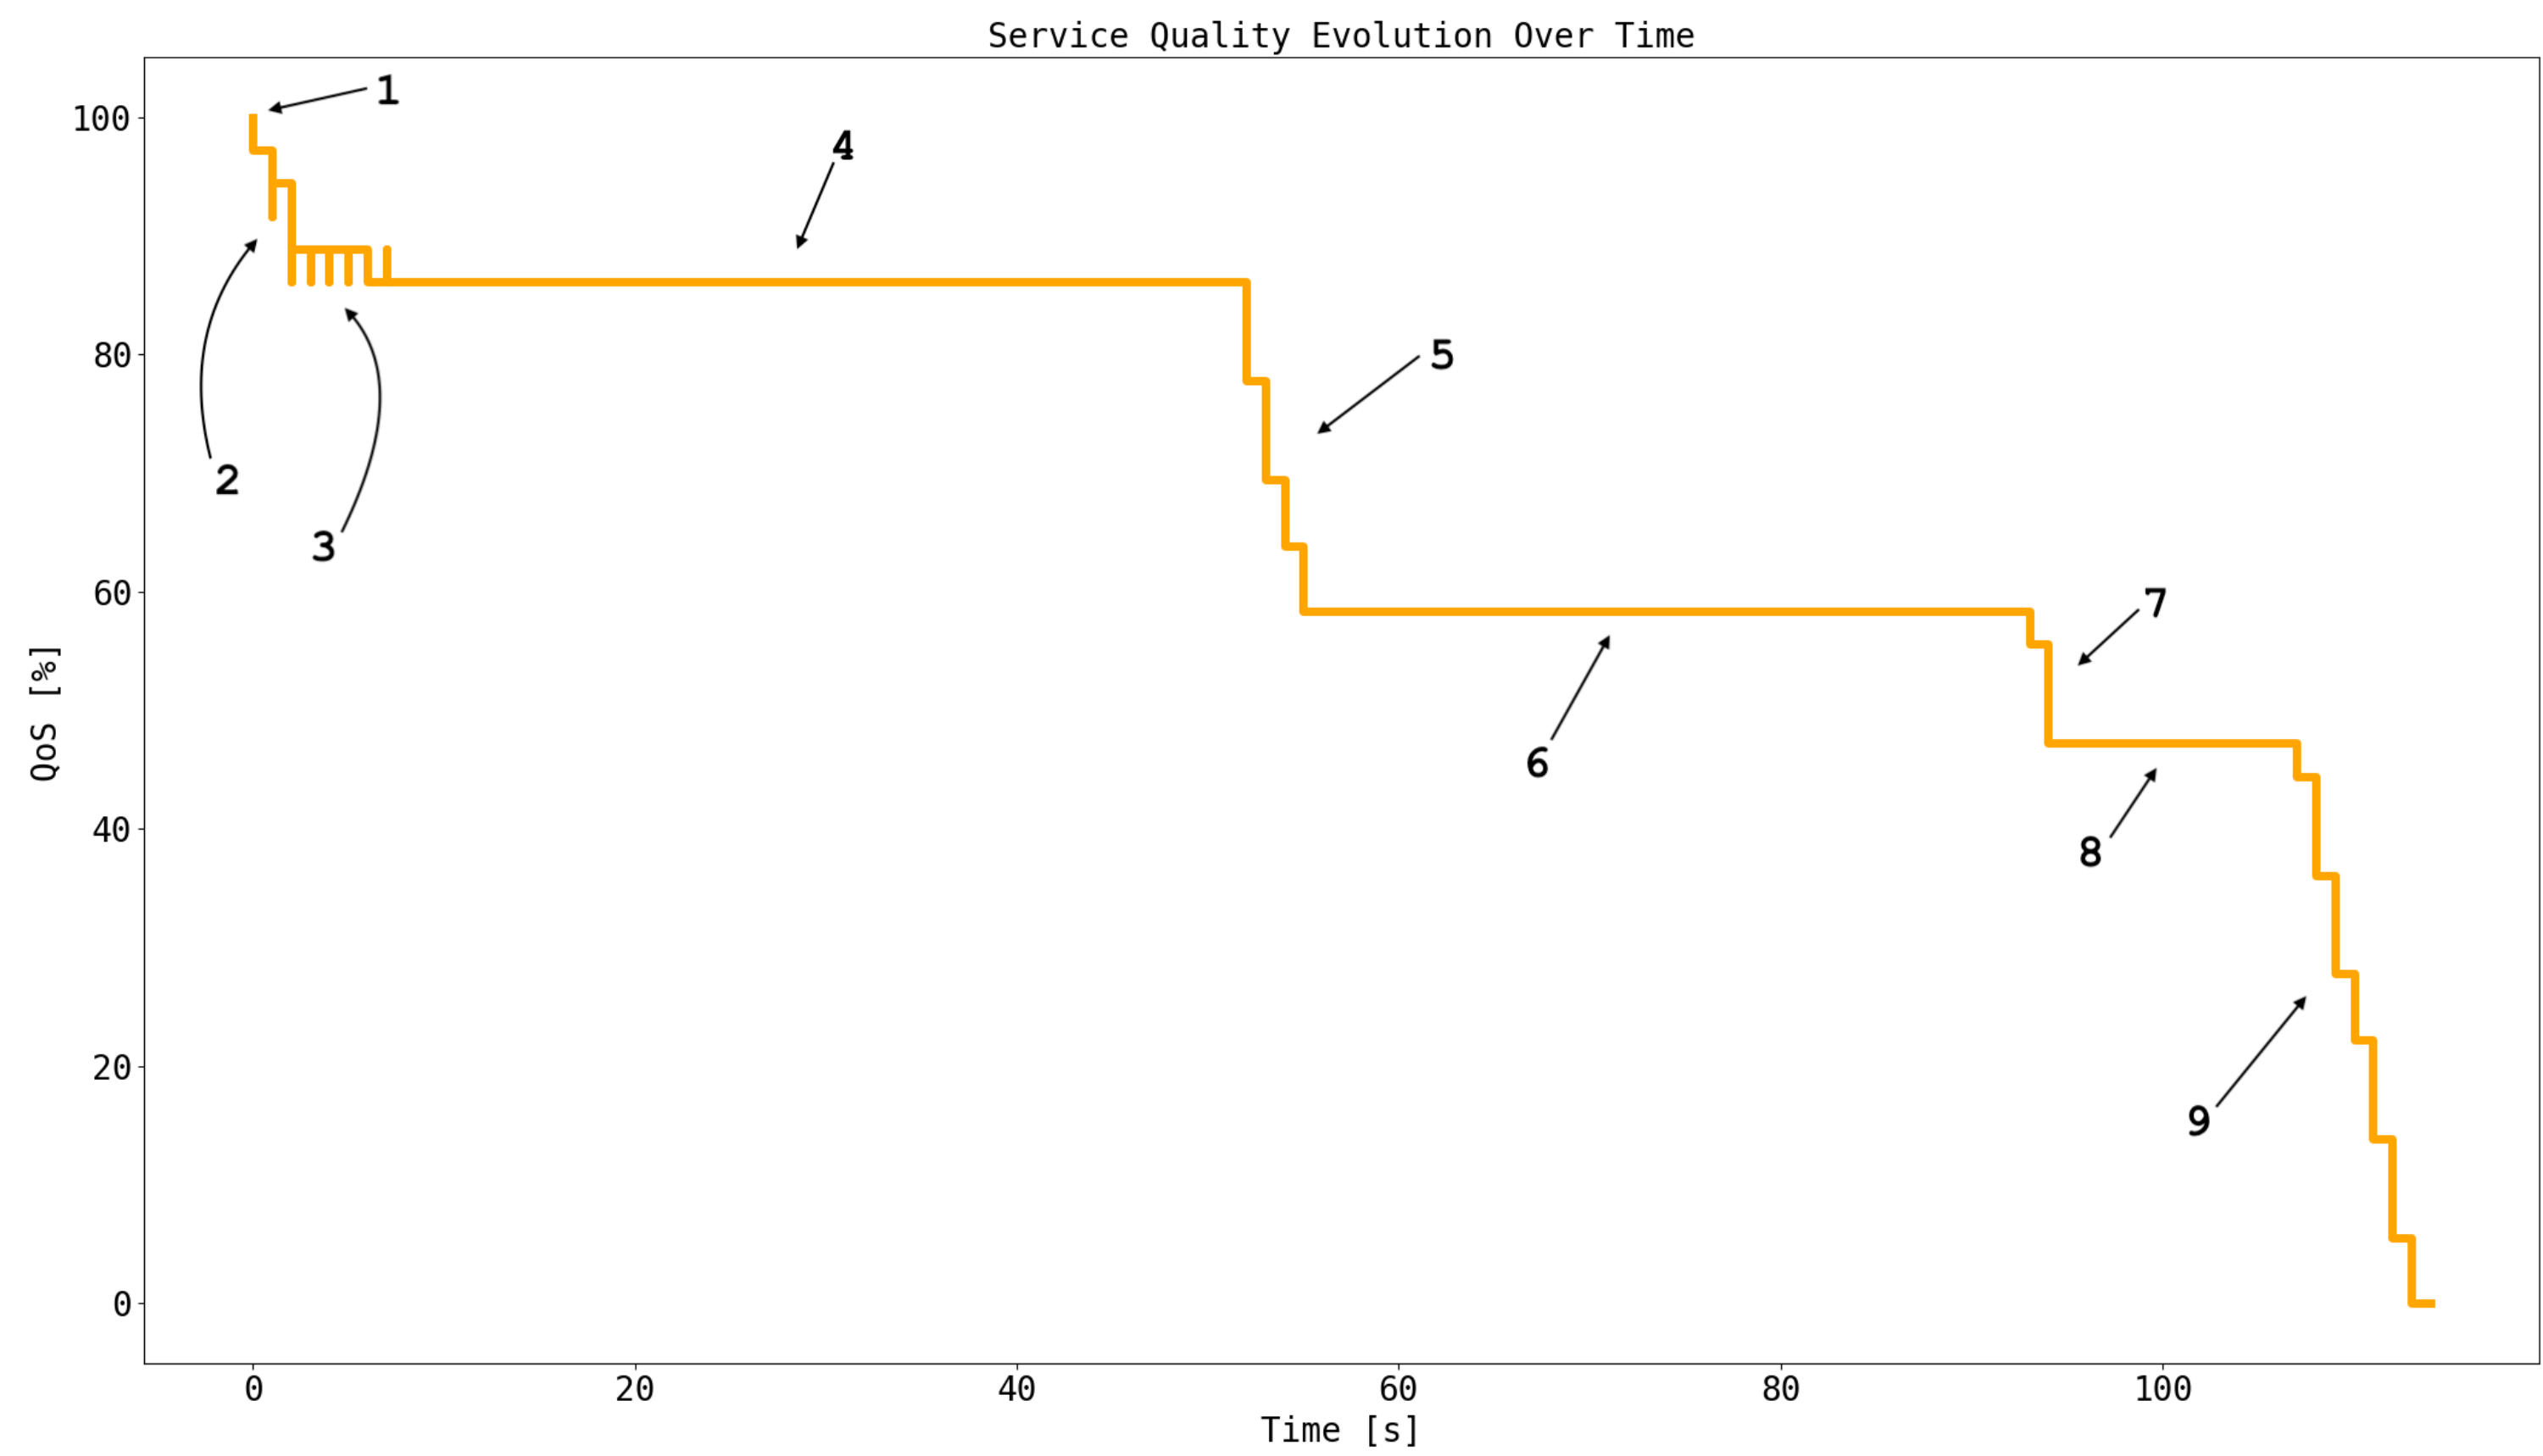
- Exploit -> weak root password (1234).
- Attack:
 - Bash script.
 - Logs into victims remotely through ssh.
 - Recursively copies itself to machines.
 - Base case -> no more subnets to attack.
 - Its progress can be known *a priori*.




QoS Evolution on a Static Network



QoS Evolution on a Dynamic Network



Comparison of QoS Evolution



Closing Thoughts and Future Work

- Huge challenge.
- Learn about many different technologies and how to mesh them.
- Gained deeper understanding of how networking works “behind the scenes”.
- Learned how to manage our time and plan more efficiently.
- Possible future use:
 - Aide professors in the realm of networking.
 - Provide a powerful and lightweight alternative to VMs for network emulation.

Thank You for Your Time and Attention

Questions?

